

QUANTIFICATIONALISM

ANTONIO MARIA CLEANI

ABSTRACT. Some text. Word count:

INTRODUCTION

Philosophers disagree about whether tense is real.¹ What exactly that means is a point of contention. One familiar piece of technology used to articulate this disagreement is the distinction between *temporary* and *eternal* propositions.² A *temporary* proposition is one that is neither always true nor always false: the proposition that I am sitting is true now, but will be false later. An *eternal* proposition, on the other hand, is one that has a fixed truth value across all times: the proposition that I am sitting at 11:00 am on February 12, 2026 is true now, has always been true, and will always be true. At a first approximation, to be a realist about tense is to hold that there are temporary propositions, while to be an anti-realist about tense is to hold that all propositions are eternal.³

There is a parallel disagreement about whether modality is real.⁴ We can elucidate it in a similar way. A *contingent* proposition is one that is neither necessarily true nor necessarily false. The proposition that Timothy Williamson is a contingentist is false, but could have been true. A *non-contingent* proposition, on the other hand, is one that has a fixed truth value across all possible worlds: the proposition that Timothy Williamson is actually a contingentist is false and could not have been true. To be a realist about modality is to hold that there are contingent propositions, while to be an anti-realist about modality is to hold that all propositions are non-contingent.⁵

Draft of October 4, 2026.

¹The debate is often framed as the question whether we should “take tense seriously.” Contemporary discussion is usually traced back to McTaggart (1908). Prior (1957, 1967) is by far the most influential realist; Gale (1968) and Fine (2006a) are other notable mentions. Prominent anti-realists include Russell (1915); Goodman (1951); Quine (1960); Mellor (1981); Lewis (1986); Sider (2001). The debate on the reality of tense is sometimes distinguished from that over the present is “metaphysically privileged” (Zimmerman, 2005; Fine, 2005), though not always (Sider, 2001, Ch. 1).

²See especially Richard (1981); Brogaard (2006); Bacon (2018); Dorr and Goodman (2020).

³Not all the anti-realists mentioned earlier count as anti-realists in my sense. Sider (2001), for example, thinks tense should be simply paraphrased away, so it is not clear whether he would accept the claim that every proposition is always true.

⁴The realist side features again Prior (1957); Prior and Fine (1977), see also Yagisawa (2010)—especially his notion of “modal tense”—and Williamson (2013b). For the anti-realist side see Spinoza (1985); Quine (1953); Schaffer (2012).

⁵As before, it is not clear whether Quine (1953) is an anti-realist in this sense, though Spinoza (1985) and Schaffer (2012) arguably do.

In this paper, I introduce another debate of this sort, which has been all but ignored in extant literature, and explore the realist side of it.⁶ The debate concerns quantification. I argue that there is sense to be made of a distinction between *domain-variable* and *domain-invariable* propositions. This distinction is to quantification what the distinction between temporary and eternal propositions is to temporal discourse. A domain-variable proposition is one whose truth value varies across *domains of quantification*: the proposition that there are no even prime numbers is false, but true *with the exception of 2*. A domain-invariable proposition, on the other hand, is one that has a fixed truth value across all domains of quantification: the proposition that except for 2, there are no even prime numbers is true, and indeed true across all domains of quantification.

A realist about domain-variability (for lack of a better term) is someone who holds that there are domain-variable propositions, while an anti-realist about domain-variability is someone who holds that all propositions are domain-invariable. I precisely formulate a particular version of realism about domain-variability, which I call *quantificationalism*, and explore the metaphysical picture that comes with it.

I argue that quantificationalism should be taken seriously because it is philosophically promising. It has a wide range of applications in topics as diverse as metaontology, the metaphysics of modality, predicativity, and (more tentatively) potentialism in the philosophy of mathematics. At the same time, it is not without costs. While precisely formulating versions of realism about tense, modality, and the first-person is relatively straightforward, doing the same for quantificationalism requires considerably more care, and naive attempts lead to paradox. I show that, nonetheless, this can be done. We can understand the notion of domain-occasional truth using the apparatus of *metaphysical substitutions* (Bacon, 2019, 2020). The resulting framework is revisionary but, I argue, ultimately illuminating.

Here is the plan for the paper. Section 1 articulates a minimal version of realism about tense, which serves as a model for the formulation of quantificationalism in Section 2. Section 3 sketches a logical framework for quantificationalism and clarifies key pieces of its ideology. Section 4 outlines some of its applications, and Section 5 develops one of them in more detail. Section 6 turns to the question of consistency, describing the changes to the framework of Section 3 required to ensure it, and exploring the resulting metaphysical picture. Section A presents the technical details of the framework in compressed form; an extended treatment can be found in the companion paper [REDACTED].

1. WARM UP: TEMPORALISM

I begin by giving a more careful articulation of quantificationalism. It will be helpful to proceed by analogy: I will first articulate a precise version of realism about tense I call *temporalism*, then introduce quantificationalism as a version of realism about domain specification modeled after temporalism.

I have formulated realism about tense as the claim that there are temporary propositions—propositions that are neither always true nor always false. In this statement, I use quantification over propositions as a mere shorthand for higher-order quantification in sentence position. A realist about tense, in my sense, is someone who accepts

⁶The lone exception appears to be Stojanovic (2012), though her framing of the debate differs significantly from mine.

Temporalism: For some p , it is neither always the case that p nor always the case that not p .⁷

The label ‘temporalism’ has been used in the literature to refer to different, and generally stronger ways of articulating realism about tense. I reserve the label for the thesis just stated above. In the remainder of this section, I will clarify the content of temporalism by comparing it with these other views.

To begin with, we should distinguish temporalism from what we may call *first-order* tense realism. This view articulates tense realism in terms of first-order quantification over propositions, a special sort of abstract objects satisfying a certain theoretical role.⁸ I find first-order tense realism an unsatisfactory formulation of tense realism, for two main reasons.

First, it threatens to make the debate over the reality of tense a merely verbal one, since it is resolved differently depending on what we take propositions to be. Propositions have been taken to serve both as the objects of attitudes like belief and desire, and as the semantic content of sentences. But it is not obvious that there is a single kind of thing that can play both these roles. Notably, [Lewis \(1980\)](#) argues the semantic contents of sentences cannot be the same as the objects of belief, on the basis that the former but not the latter vary in truth value across times and worlds. If Lewis is right, we cannot take propositions to play both roles. We must make a choice, and the truth of realism about tense depends on which choice we make.

Second, one might worry that this formulation saddles realism about tense with irrelevant ontological commitments. No matter whether propositions are taken to be semantic values of sentences or objects of attitudes, most philosophers agree that propositions are abstract objects. Nominalists, who do not believe in abstract objects, do not believe in propositions, and a fortiori do not believe in temporary propositions. Thus, no nominalist can be a realist about tense.

Temporalism does better than first-order tense realism on both these fronts. In classical higher-order logic, the meaning of quantifiers in sentence position is uniquely determined by the introduction and elimination rules governing them, up to logical equivalence ([Harris, 1982](#)). Moreover, as argued at length by proponents of higher-order logic (e.g., [Bacon, 2024](#)), quantification in sentence position does not carry any ontological commitment to propositions, understood as first-order abstract objects. I thus set aside first-order tense realism, and throughout the paper use expressions like ‘some propositions’ as shorthands for higher-order quantification in sentence position. Similar conventions apply to expressions like ‘some properties,’ ‘all relations,’ and so on.

Tense realism is also sometimes understood in terms of fundamentality-theoretic notions. For example, one might understand tense realism as the claim that, *fundamentally*, there are temporary propositions.⁹ On this construal of the debate, someone might accept the existence of temporary propositions while still counting

⁷Temporalism in the sense just defined is discussed in [Bacon \(2018\)](#); [Dorr and Goodman \(2020\)](#), and it is implicit in Arthur Prior’s work.

⁸This is the version of temporalism found in [Brogaard \(2006\)](#) and denied, e.g., by [Richard \(1981\)](#).

⁹Examples include [Zimmerman \(2005\)](#); [Fine \(2005\)](#). Of course, these views can be further distinguished by what sort of fundamentality-theoretic ideology they employ.

as an anti-realist about tense. For example, they might think that all temporary propositions are *grounded* in eternal propositions, and therefore not fundamental.

As I stated, temporalism does not make any claim about the fundamentality of temporary propositions. It simply says that there are some. I chose this weaker formulation because I find fundamentality-theoretic strengthenings of temporalism unnecessary distractions at the present dialectical stage. My goal is to introduce a philosophical position, realism about domain-variability, by analogy with tense realism. Since realism about domain-variability has not been articulated before, it makes sense to analogize it with the simplest version of tense realism, which I take to be temporalism.

In a similar vein, sometimes the central questions in the debate over the reality of tense are taken to be about how tense operators like ‘always’ and ‘it was the case that’ relate to quantification over times. A realist insists that tense operators are primitive, and quantification over times is legitimate only if analyzed in terms of tense operators (Prior and Fine, 1977). On the other hand, an anti-realist insists that quantification over times is legitimate in its own right, and tense operators must be reduced to quantification over times.

How, exactly, temporalism relates to this construal of the debate depends on what we make of the notions of “reduction” and “analysis.” But there is at least one sense in which temporalism is perfectly compatible with the claim that quantification over times is legitimate in its own right. Consider:

Tense Leibnizianism: For any p , for it to be the case that, sometimes, p is for there to be a time t such that, at t , p .

If we deny that quantification over times must be reduced to the tense operators ‘sometimes’ and ‘always,’ then this appears to be a definition of ‘sometimes’ in terms of quantification over times. But as long as there is some p and some times t and t' such that, at t , p but at t' , not p , tense Leibnizianism is consistent with temporalism. Indeed, given tense Leibnizianism, a temporary proposition can be *defined* as a proposition p such that, at t , p but at t' , not p , for some times t and t' .

This point bears emphasis. Temporalism, essentially, is the thesis that the tense operators ‘sometimes’ and ‘always’ are non-trivial: they can turn a truth into a falsehood or vice versa. This claim is orthogonal to the question whether we regard quantification over times as legitimate in its own right or not. A temporalist is free to accept times as legitimate entities and use them to introduce additional, more specific tense operators that regiment the intuitive notion of truth at a time. Namely, for each time t , one introduces a further tense operator ‘at t .’¹⁰ These operators will also generally be non-trivial: for most t , a temporalist will claim there is some p such that either p but at t , not p , or the other way around. These more specific tense operators can be used to illuminate the logical behavior of the more general tense operators ‘sometimes’ and ‘always,’ as done in tense Leibnizianism, but they need not be.

¹⁰Prior (1967); Prior and Fine (1977) construct instants as propositions that are sometimes the case and strictly imply, in the sense of eternal material implication, any proposition they are sometimes true together with. In such a setting ‘at t ’ can be defined as the operator of being eternally implied by t . To Prior and Fine, this kind of quantification over times is legitimate because times themselves are construed in terms of tensed ideology. A temporalist in my sense, on the other hand, is free to embrace tense Leibnizianism while rejecting any analysis along these lines, so long as they maintain that ‘at t ’ is a sentential operator.

Philosophers who consider themselves anti-realists about tense sometimes endorse a different, *linguistic* reducibility claim. They claim that, contrary to appearance, what I have been calling tense operators are not sentential operators at all. These expressions do involve quantification over times, but the truth conditions of sentences containing them can be given using only eternal propositions.

Let me elaborate with an example. Consider the sentences:

- (1) No country has all five of Earth's climate zones.
- (2) In 1783, no country had all five of Earth's climate zones.

Now, (1) is false, because the United States does have all five of Earth's climate zones. In fact, it is the only country to ever enjoy this much climate diversity in at least half a century. But (2) is true: the western border of the United States in 1783 was the Mississippi River, so dry climates were not present in the United States at that time.

What explains the difference in truth value between (1) and (2)? A temporalist has a simple explanation: (2) is constructed by applying a non-trivial tense operator, 'in 1783', to (1). The latter expresses the same proposition when it occurs within the scope of 'in 1783.' Because 'in 1783, p ' entails 'sometimes, p ,' (1) is a temporary falsehood: false now, but true in the past. Call a temporalist that goes for an explanation along these lines a *linguistic temporalist*.¹¹

Anti-realists who deny temporalism offer an alternative explanation for the difference in truth value between (1) and (2).¹² They insist that both (1) and (2) express eternal propositions, and deny that (1)—if it expresses a proposition at all—expresses the same proposition as its embedded occurrence in (2). Here is one way this thought can be fleshed out further. Somewhere in the syntax of (1) there is a free occurrence of a covert variable that takes times as values. When one utters (1) at a time t , the context of utterance assigns the value t to this variable, and the utterance expresses the proposition that, at t , no country has all five of Earth's climate zones. When t is, say, 11:00 am on February 12, 2026, the resulting proposition is an eternal falsehood. The modifier 'in 1783' in (2) works by binding this time variable and specifying some moment in 1783 as its value. Thus, for any t , when one utters (2) at t , that utterance expresses the eternal truth that, in 1783, no country had all five of Earth's climate zones. Thus, we have an explanation of the variation in truth value between (1) and (2) that does not appeal to temporary propositions.

Call someone who gives an explanation along these lines of tensed discourse a *linguistic anti-temporalist*. Clearly, anyone who does not believe in temporary propositions must go for some version of linguistic anti-temporalism. However, I claim, temporalism, on its intended reading, is compatible with both linguistic temporalism and linguistic anti-temporalism.

This point might seem surprising, so allow me to elaborate. Suppose you like the linguistic anti-temporalist explanation of (1) and (2). You think that, as a matter of natural language syntax and semantics, (1) and its embedded occurrence in (2) express different propositions, and that the modifier 'in 1783' is not a sentential

¹¹Prior (1957, 1967) is a linguistic temporalist in my sense. Brogaard (2006) endorses a similar linguistic thesis, though formulated in first-order terms. See also Kaplan (1989) and the subsequent literature on "operator arguments" (Yli-Vakkuri et al., 2019).

¹²This is the standard story told by all previously cited tense anti-realists, though in first-order terms.

operator at all. Even so, the following question can be raised: is there a way to *introduce* a sentential operator that behaves like ‘in 1783’ in the sense that it can be used to express the same propositions we express in English by prefixing sentences with ‘in 1783’?

More precisely, is it true that for any time t , there is an operator $\text{at}(t)$ such that, for any sentence φ , there is a proposition p with the following properties:

- The proposition expressed by φ when uttered at t is identical to $\text{at}(t)(p)$, and
- The proposition expressed by ‘in 1783, φ ’ when uttered at t is identical to $\text{at}(1783)(p)$?

Someone who answers ‘yes’ to this question counts as a temporalist, even though they believe tensed discourse in natural language is best accounted for without using temporary propositions.

In view of this, we should reformulate temporalism slightly:

Temporalism: For some p , it is neither always* the case that p nor always* the case that not p .

Here, ‘it is always* the case that’ is a quasi-technical expression that stands to the English expression ‘it is always the case that’ in the same way that ‘ $\text{at}(1783)$ ’ stands to the English expression ‘in 1783.’ A linguistic anti-temporalist might not think that ‘it is always the case that’ is itself a sentential operator, but for independent reasons might posit a sentential operator ‘it is always* the case that’ that can be used to express the same propositions we express in English by prefixing sentences with ‘it is always the case that.’ On these grounds, they count as a temporalist in my sense.

As I hope the discussion above illustrates, temporalism is, at its core, a thesis about *granularity*. Temporalism asserts that reality is fine-grained enough for the notion of temporary truth to make sense. Extremely coarse-grained pictures of reality falsify temporalism. For example, assume propositions are individuated by material equivalence: two propositions are the same iff they have the same truth value. It follows that there are only two propositions, truth and falsity. In this case, temporalism must be false. For surely logical truths are always true, and contradictions are always false. So, all true propositions are always true, and all false propositions are always false. But any proposition is identical, on this picture, with either a logical truth or a contradiction. There simply are not enough propositions around for us to introduce operator expressions that faithfully regiment tensed discourse.

2. FROM TEMPORALISM TO QUANTIFICATIONALISM

At a first approximation, *domain-variability* is the apparent dependence of truth and falsity on what is included in the domain of quantification. It is false that all prime numbers are odd, but *with the exception of 2*, all prime numbers are indeed odd. Some mammals lay eggs, yet it is false that, *in addition to the Platypus and the Echidna*, some (other) mammals lay eggs.

A realist about domain-variability is someone who holds that there are domain-variable propositions, which change their truth values across domains of quantification much the same way as temporary propositions change their truth values across

times. My preferred version of realism about domain-variability is *quantificationalism*, which I formulate as follows:

Quantificationalism: For some p , it is neither domain-invariably the case that p , nor domain-invariably the case that not p .

As in temporalism, quantificationalism cashes out quantification over propositions as higher-order quantification in sentence position. The expression ‘it is domain-invariably the case that’ is a theoretical term stipulated to occupy the position of a sentential operator, and expressing generality with respect to domains of quantification the way that ‘it is always the case that’ expresses generality with respect to time. I will introduce the term ‘it is domain-occasionally the case that’ as an abbreviation for the operator ‘it is not domain-invariably the case that not,’ which analogizes ‘it is sometimes the case that.’ While certainly less familiar than their temporal analogue, I hope the content of these operators will get clearer as we unpack quantificationalism further.

It is important not to conflate quantificationalism with a number of theses in its vicinity. Here, the lessons of the previous section carry over. Quantificationalism is a higher-order generalization, and thus must be distinguished from the claim that there are domain-variable propositions, understood as a special sort of abstract objects satisfying a certain theoretical role.¹³ Quantificationalism makes no claim about the fundamentality of domain-variable propositions, and is independent of the idea that quantification over domains (whatever that may be) is legitimate in its own right.

Natural language contains expressions that appear to operate on domains of quantification, and we need to be careful not to conflate quantificationalism with a thesis about the syntax and semantics of these expressions. It is certainly true that one *might* be drawn to quantificationalism as a framework for providing a syntactic and semantic analysis of these expressions. Consider:

- (3) No country has all five of Earth’s climate zones.
- (4) With the exception of the United States, no country has all five of Earth’s climate zones.

We noted earlier that the United States is the only country to have all five of Earth’s climate zones. Thus, (3) is false, while (4) is true. A realist about domain-variability can explain this difference in truth value mirroring the linguistic temporalist explanation of the difference in truth value between (1) and (2). They might say that (4) is constructed by applying a non-trivial operator expression, ‘with the exception of the United States,’ to (3). The latter expresses the same proposition when it occurs within the scope of ‘with the exception of the United States.’ Because ‘with the exception of the United States, p ’ entails ‘domain-occasionally, p ,’ (3) is a domain-occasional falsehood: false, but true relative to the domain of things other than the United States. A quantificationalist that goes for an explanation along these lines is a *linguistic quantificationalist*.¹⁴

Yet, as before, this sort of account is by no means forced upon the quantificationalist. Quantificationalists, who might have independent reasons to believe in domain-invariable propositions, are free to deny linguistic quantificationalism, and

¹³Stojanovic (2012) defends a position of the latter sort. Stojanovic calls *domain sensitivity* roughly what I call domain-variability, though hers is a first-order notion.

¹⁴Stojanovic (2012) endorses a first-order version of linguistic quantificationalism.

opt to explain the difference in truth value between (3) and (4) purely in terms of domain-invariable propositions.

Here is how this alternative explanation might go. For whatever reason, when one utters (3) in a context where the salient domain of quantification includes the United States, the utterance expresses a domain-invariable falsehood: the proposition that, *in the contextually salient domain d* that includes the United States, no country has all five of Earth’s climate zones. If, in the same context, one utters (4), the utterance expresses a domain-invariable truth: the proposition that, in the domain that results by removing the United States from d , there is no country having all five of Earth’s climate zones. Perhaps this happens because the syntax of (3) contains covert occurrences of free variables ranging over domains of quantification.¹⁵

To the extent they wish to tie genuine domain-variability to natural language, a quantificationalist who denies linguistic quantificationalism might insist that, even though ‘with the exception of the United States’ is not itself a sentential operator, it is possible for us to *introduce* a sentential operator that we can use to express the same propositions we express in English by prefixing sentences with ‘with the exception of the United States.’ More precisely, this sort of quantificationalist will insist that, given any domain D , there is an operator $\text{At}(D)$ such that, for any sentence φ , there is a proposition p with the following properties:

- The proposition expressed by φ when uttered in a context where the salient domain of quantification is D is identical to $\text{At}(D)(p)$, and
- The proposition expressed by ‘with the exception of the United States, φ ’ when uttered in a context where the salient domain of quantification is D , is identical to $\text{At}(D \setminus \{\text{United States}\})(p)$, where $D \setminus \{\text{United States}\}$ is the result of removing the United States from D .

At the very least, the logic of this operator $\text{At}(D)$ will include the claim that ‘ $\text{At}(D)(p)$ ’ entails ‘it is domain-occasionally the case that p .’ So, on a view of this sort, quantificationalism can be inferred from (3) and (4), even though (3) itself expresses a domain-invariable falsehood.

Just like temporalism, then, quantificationalism is, at its core, a thesis about granularity. It asserts that reality is fine-grained enough for the notion of domain-invariable truth to make sense. Extremely coarse-grained pictures of reality, such as the previously considered Fregean view that propositions are individuated by material equivalence, falsify quantificationalism. Indeed, presumably logical truths are domain-invariably true, and contradictions are domain-invariably false. So, all true propositions are domain-invariably true, and all false propositions are domain-invariably false. There simply are not enough propositions around for us to introduce operator expressions that faithfully regiment discourse involving domain-manipulation.

3. GETTING MORE PRECISE

Quantificationalism is only as clear as the notion of domain-occasional truth. While, I hope, a rough idea of what this unfamiliar notion is has emerged from

¹⁵This is the dominant approach for dealing with cases like (1) and (4), found, e.g., in Westerstähl (1985); Stanley (2002). Stanley and Szabó (2000) contains relevant discussion of contextual quantifier domain restriction, though it does not deal with explicit domain restriction by a phrase like ‘with the exception of the U.S.’

discussion so far, the best way to properly clarify it is to describe its logic. That is the goal of the present section.

To that end, let me continue the analogy between temporary truth and domain-occasional truth. One approach to developing the logic of temporary truth is to connect the notion of temporary truth to the notion of truth at a time, as tense Leibnizianism does, then read off the logic of temporary truth from the logic of truth at a time. This is by no means the only way to develop the logic of temporary truth. But it is a natural one that, as I argued in Section 1, is open to the temporalist as long as they cash out truth at a time in terms of sentential operators. I take a similar approach to the logic of domain-occasional truth. I introduce a notion of truth at a *particular* domain, connected to domain-occasional truth roughly the way that truth at a time is connected to temporary truth.

The goal of this section is to outline a logical framework for quantificationalism based on this idea at a level of detail suitable for a general philosophical audience.¹⁶ I must make clear at the outset that the version of the framework outlined here is *not* intended to be definitive. In Section 3.1, I will make a few assumptions about syntax (in particular, λ -abstraction and quantification) that I think quantificationalists should ultimately reject. I make these assumptions with the goal of starting off from a place that will feel more familiar to readers with some exposure to the higher-order metaphysics literature. I will flag these assumptions using footnotes.

The logically minded reader may notice a serious problem with the system outlined in this section. I deal with that problem in Section 6.1, and show how it can be avoided by dropping the aforementioned syntactic assumptions.¹⁷ Avoiding paradox, however, is not the sole motivation for syntactic revisionism. In Section 3.3, I propose to understand the notion of truth at a domain in terms of *metaphysical substitutions* (Bacon, 2019, 2020). Syntactic revisionism is independently motivated by the goal of perspicuously representing this interpretation of truth at a domain syntactically.

3.1. A Higher-Order Language. We will work in a simply typed higher-order language. In this language, terms come categorized into *types*. There is a basic type e of singular terms, and a basic type t of formulas. There are functional types. For example, we have a type $e \rightarrow t$ of first-order monadic predicates: expressions that combine with a singular term to yield a formula. We have a type $t \rightarrow t$ of sentential operators: expressions that combine with a formula to yield a formula. More generally, whenever σ and τ are types, there is a functional type $\sigma \rightarrow \tau$ of expressions that combine with an expression of type σ to yield an expression of type τ . Since ours is a simply typed system, no typed term can combine with more than one typed expression at once. A first-order binary relation is treated as an expression that combines with a singular term to yield a first-order monadic predicate. For example, to construct the sentence ‘Alice loves Bob’ the binary predicate ‘... loves ...’ first combines with ‘Alice’ to yield the first-order monadic

¹⁶An extended technical treatment can be found in [REDACTED].

¹⁷I will flag the relevant syntactic assumptions using footnotes briefly describing their alternatives. The reader whose main interest is to get a sense of the possible applications of quantificationalism should feel free to ignore these for the time being. The reader more interested in foundational questions will find pointers to relevant discussion in Section 6.1 in these footnotes. Section 6.1 can be read before Sections 4 and 5, but should be read after the current section.

predicate ‘Alice loves ...,’ which then combines with ‘Bob’ to yield the desired sentence.

What makes ours a “higher-order” language is the ability to quantify in every grammatical position, not just that of singular names. The way this is usually done in the literature, and that I shall conform to for the time being, is to introduce a logical constant $\forall_\sigma$ of type $(\sigma \rightarrow t) \rightarrow t$ for each type σ .¹⁸ To assert that everything is F , where F has type $\sigma \rightarrow t$ we write $\forall_\sigma F$. This approach is reminiscent of the Fregean idea that quantifiers are properties of properties: to say that everything is F is to say something like F is *universal*.

In addition to higher-order quantifiers, I assume our constants include the usual Boolean connectives (negation, $\neg$, has type $t \rightarrow t$ and all binary Boolean connectives have type $t \rightarrow t \rightarrow t$) as well as a constant $\equiv_\sigma$ of type $\sigma \rightarrow \sigma \rightarrow t$ for each type σ . For now, we can proceed under the assumption that these constants express higher-order identity relations at each type.¹⁹

Constants, together with an infinite supply of variables for each type, constitute the basic vocabulary of our language. From these we build complex terms using four term formation rules, two of which are standard and two of which are distinctly quantificationalist.

The first is *application*. This is the term formation rule I have been alluding to when speaking of terms *combining* with one another to describe the type hierarchy. In general, whenever we have a term M of type $\sigma \rightarrow \tau$ and a term N of type σ , the result of *applying* M to N , written MN , is a term of type τ .

The second rule is *λ -abstraction*. If M is a term of type τ and x is a variable of type σ , then the expression $\lambda x.M$ is a term of type $\sigma \rightarrow \tau$. Note that $\lambda x.M$ is well-formed regardless of where x occurs in M , or whether it occurs at all.²⁰ The meaning of λ -terms is fixed by two main stipulations:

$$\begin{aligned} (\lambda x.M)N & \text{ is synonymous with } M[N/x] & (\beta) \\ (\lambda x.Mx) & \text{ is synonymous with } M, & (\eta) \end{aligned}$$

where $M[N/x]$ denotes the result of substituting N for x in M . When M and N are connected by chains of synonymy of either kind above, they are called *$\beta\eta$ -equivalent*. The background logic allows *$\beta\eta$ -equivalent* terms to be substituted for one another *salva veritate* in all contexts.

Next, we have the two quantificationalist term formation rules. First, if P is a formula, we can form a new formula, $\blacksquare P$. This is how I regiment the claim ‘it is domain-invariably the case that P .’ So, in the present framework, quantificationalism can be stated as follows:

Quantificationalism: $\exists p(\blacklozenge p \wedge \blacklozenge \neg p)$.

¹⁸The alternative is to treat quantifiers as *syncategorematic expressions*, in a sense to be explained shortly. Roughly, this is the way quantifiers are usually treated in first-order logic: they are introduced via a rule specifying how to construct new terms from old, not as self-standing expressions. In the final version of the framework, I will opt for a syncategorematic treatment.

¹⁹In the final version of the framework, these constants will end up expressing relations weaker than identity that I call relations of *applicative indiscernibility*, in terms of which higher-order identity relations themselves can be defined. See Section 6.3.

²⁰In the final version of the framework, I will impose restrictions on which variables can be abstracted. See Section 6.2.

Second, if F is a predicate of type $\sigma \rightarrow t$ and M is a term of type τ , we can form a new term, $\text{At}(F)(M)$, also of type τ . An expression of the form $\text{At}(F)$ is called a *domain specifier*.

A domain specifier $\text{At}(F)$ is an expression whose semantic function is to specify precisely the F s as the domain of quantification within its scope. Intendedly, this is analogous to the way that expressions like ‘in 1783’—when thought of as sentential operators—force sentences within their scope to be evaluated with respect to the time 1783. In this framework, a sentence like ‘with the exception of the United States, no country has all five of Earth’s climate zones’ is regimented as something like

$$\text{At}(\lambda x.x \neq \text{United States})(\neg \exists x(\text{Country}(x) \wedge \text{HasAllFiveClimateZones}(x))).$$

Henceforth, I will adopt the convention of pronouncing a domain specifier $\text{At}(F)$ as ‘among the F s.’ I want to stress, however, that this is really just a pronunciation convention, not a substantive hypothesis about what expressions of the form ‘among the F s’ mean in English.

Domain specifiers regiment the informal notion of truth at a domain. Domains, in my theory, are identified with higher-order properties, and any such property corresponds to a domain.²¹ For a proposition P to be true at the domain consisting of the F s is simply for $\text{At}(F)(P)$ to be true *simpliciter*. Domain specifiers relate to the domain-invariable truth operator $\blacksquare$ in roughly the same way tense nominals relate to the tense operator ‘always.’ In one direction, this means that if $\blacksquare P$ is true, then so is $\text{At}(F)(P)$, for any domain specifier $\text{At}(F)$. In the other direction, the intuitive idea is that if $\text{At}(F)(P)$ is true for all domain specifiers, then $\blacksquare P$ is true, though some care is needed when formulating this claim precisely.²²

Expressions like constants, variables, and typed terms that can be iteratively built out of constants and variables through the term formation rules just reviewed are called *categorematic*. Their semantic function is to stand for higher-order entities. Domain specifiers and $\blacksquare$ are *syncategorematic* expressions. They are devices for syntactically notating ways in which typed terms themselves are built. Their semantic function is not to stand for higher-order entities, but to describe ways in which higher-order entities can be combined.

All higher-order languages need some syncategorematic expressions. For all higher-order languages need term formation rules and term formation rules must be notated somehow, so as to allow us to uniquely parse the syntactic structure of a term. Standard higher-order languages typically have two such expressions: application and λ . I am here using the term ‘expression’ in a loose sense that counts term concatenation—the device we use to notate that a term has been built by applying one to another—as an expression. First-order languages treat quantifiers and Boolean connectives as syncategorematic expressions rather than constants, and propositional modal languages treat modal operators the same way.

²¹Officially, domain specifiers take as their first argument a *finite sequence* of properties of distinct types, so that the domain can be specified at several types at once (Section A). I ignore this complication in the main text.

²²The main problem is that At can take a predicate for any type in its first argument place, and our higher-order language lacks the ability to quantify over properties of any type at once. We can get around this complication by using inference rules: roughly, if we can prove $P \rightarrow \text{At}(F)(Q)$ for an *arbitrary* domain specifier $\text{At}(F)$, then we can also prove $P \rightarrow \blacksquare Q$.

I take the distinction between categorematic and syncategorematic expressions to be more than merely notational. This attitude has historical precedent, especially with regards to logical vocabulary. The idea is that logical vocabulary should be treated syncategorematically (the way quantifiers and Boolean connectives are in first-order logic) to reflect the fact that it is not meaningful in virtue of contributing content or constituents to propositions, but rather by describing the “form” of propositions, or the way their constituents are arranged. This idea stretches back as far as [Buridan \(1976, I.7.2\)](#) and other medieval logicians (see [Kretzmann, 1982](#)), and is perhaps most famously championed by the early logical atomists, e.g., [Wittgenstein \(1921, §4.0312, §5.4611\)](#) [CHECK] and [Russell \(1940, pp. 39–40\)](#).

While I do not treat all logical vocabulary as syncategorematic here, I am sympathetic to the general idea of treating expressions whose semantic function is to describe the way in which entities can be meaningfully combined syncategorematically. I treat quantificationalist vocabulary this way because this is precisely its role. Domain specifiers, and by extension ■, describe a way of putting higher-order entities together that differs from and should not be reduced to application. In [Section 3.3](#) I elaborate on what it consists in.

3.2. Free Logic. We are trying to study expressions capable of manipulating domains of quantification, so it is important to get a handle on what the logic of quantification itself must be like in order for this to be possible. The most important point to that effect is that quantifiers whose domain is capable of being so manipulated must obey a *free* rather than *classical* logic.

Classical quantificational logic validates the principle of unrestricted existential generalization:

$$Fx \rightarrow \exists x Fx. \quad (\text{EG})$$

From the claim that Gavin Newsom is tall, I can infer that something is tall. The distinguishing feature of free logics is that they admit exceptions to this principle.

One reason to adopt a free logic is a desire to reason with empty terms. Using [\(EG\)](#), it is possible to derive every instance of the schema

$$\exists x(x \equiv t).$$

If we take t to be synonymous with the name ‘Sherlock Holmes,’ the corresponding instance of the above says that there is something identical to Sherlock Holmes, a claim you should resist if you think that ‘Sherlock Holmes’ does not refer to anything.

Another route to rejecting [\(EG\)](#), which is more pertinent here, comes from the idea that existence is metaphysically contingent: some things that do exist might have failed to exist. The idea of metaphysically contingent existence is ruled out by [\(EG\)](#) together with standard modal logic, from which one can derive:

Necessitism: Necessarily, everything necessarily exists.

$$\Box \forall x \Box \exists y (x \equiv y)$$

The modal logic required to show this is just the rule of necessitation for $\Box$: if P follows from no assumptions, then so does $\Box P$.²³

Nothing in the underlying argument hinges on $\Box$ being metaphysical necessity. In exactly the same way, we can show that if [\(EG\)](#) is a logical truth and ■ obeys

²³The derivation, and more discussion of necessitism, can be found in [Williamson \(2013a\)](#).

the rule of necessitation, then domain-occasional existence is impossible: anything that in fact exists domain-invariably does.

Domain-invariabilism: It is domain-invariably the case that everything is domain-invariably identical to something.

$$\blacksquare \forall x \blacksquare \exists y (x \equiv y)$$

A quantificationalist cannot accept domain-invariabilism. The very idea driving quantificationalism is that we can evaluate propositions for truth or falsity relative to non-current domains of quantification, some of which are smaller than the current one. But domain-invariabilism says exactly that domains of quantification cannot shrink.

Quantificationalists, therefore, must either adopt a free logic or deny that $\blacksquare$ obeys the rule of necessitation. I will take the first option. But I should stress that I do not think the dispute over which of these two alternatives we should take is a deep one. We could insist that every instance of (EG) is a logical truth but reject the rule of necessitation for $\blacksquare$ on the grounds that logical truths need not be domain-invariably true. Or we could insist that logical truths are indeed domain-invariably true and consequently reject (EG). The latter option, note, is perfectly compatible with maintaining that every instance of (EG) is *in fact* true—as it turns out, we are in fact quantifying over an absolutely unrestricted domain—just not domain-invariably so. This appears to me to be a merely verbal disagreement over how to use the notion of logical truth.

To theorize in a free logic is to insist that our primitive quantifier $\forall$ does not have a classical logic, in the sense that it does not satisfy (EG) within the scope of $\blacksquare$. Doing so does not foreclose the possibility of introducing some other quantifier, distinct from $\forall$, which does have a classical logic. Views that theorize using a combination of a free quantifier and a classical quantifier are in fact quite common in the literature. *Contingentism*—the negation of necessitism—is a prominent example.²⁴ On pain of inconsistency, contingentists must either embrace a free logic or deny the rule of necessitation for metaphysical necessity; in practice, they tend to take the former option. Contingentists typically claim that existence is contingent not only in the sense that some things that do exist might have failed to exist, but also in the sense that there could have been things that do not in fact exist. For example, they might claim that while Wittgenstein could have had a first child, in fact there is nothing that could have been Wittgenstein’s first child. When making generalizations, contingentists often help themselves to an “outer” quantifier whose logic is assumed to be classical, and whose range includes things such as Wittgenstein’s merely possible first child (Fritz and Goodman, 2016).

Like the contingentist, I deny that $\forall$ itself has a classical logic, but I will make use of a distinct quantifier whose logic *is* classical. In fact, I will show that such a quantifier can be *defined* from $\forall$ using domain specifiers. What I take the discussion of domain-invariabilism above to show is merely that no quantifier whose domain can be manipulated by $\blacksquare$ and domain specifiers in the ways I shall describe can have a classical logic. Being “flexible” this way and having a classical logic are two interesting but incompatible properties of quantifiers that can be useful for different purposes. It can thus be desirable for a framework to be able to express both.

²⁴See Williamson (2013b); Fritz and Goodman (2016) for discussion. We find a similar approach among actualists, who believe that everything is actual. See especially (Fine, 1977, 2016).

The system of free logic I shall use here is axiomatized using infinitely many constants $E!_\sigma$, each of type $\sigma \rightarrow t$, for each type σ . These are called *existence predicates*. (EG) says that if P , then there is something such that P . The present system replaces that with a weaker claim: if P is true *and x exists*, then there is something such that P :

$$Fx \rightarrow (E!x \rightarrow \exists x Fx). \quad (E!EG)$$

The notions of existence expressed by these predicates are quantificational. The background logic proves that $E!_\sigma$ and the predicate *is something*, $\lambda x^\sigma. \exists y(x \equiv y)$, are equivalent. Because of this, the addition of existence predicates does not correspond to an increase in ideological complexity. We could in principle let the symbol $E!_\sigma$ be just an abbreviation of the latter and not much would change. The reason why I take them as primitive is to emphasize the central role they play in clarifying the semantic function of domain specifiers, which I discuss in the next section.

A property is called *existence-entailing* when having it entails existing. I will be as liberal as possible regarding which properties are existence-entailing. Particular properties that can be logically built from existence or quantifiers are guaranteed to be existence-entailing: for example, *being an existing philosopher* ($\lambda x. E!x \wedge Px$), *being someone who loves themselves* ($\lambda x. \exists y(Lxy \wedge x \equiv y)$) or existence itself. But the logic is consistent with a wide range of hypotheses about whether properties that cannot be so built are existence-entailing. For example, it would be perfectly consistent to maintain that Sherlock Holmes is a fictional detective, or that he lives at 221B Baker Street, while denying that Sherlock Holmes exists.²⁵

3.3. Metaphysical Substitutions. I said that domain specifiers function by specifying a particular property as the domain of quantification within their scope. I now elaborate on what that means.

Recall that I think of domains of quantification simply as higher-order properties. Each higher-order property F represents a domain consisting precisely of the F s. At each type, there is a special domain, *the current domain*, represented by the property of *existing* (at that type). Part of what it means to say that a domain specifier, $At(F)$, specifies the F s as the domain of quantification within its scope, can be spelled out in terms of how $At(F)$ acts on the current domain (of the same type):

$$At(F^{\sigma \rightarrow t})(E!_\sigma) \equiv F^{\sigma \rightarrow t}. \quad (Gen)$$

In words: *to exist among the F s just is to be F .*

It is important to draw a distinction between domain *specification* and domain *restriction*. In a free logical setting, it is perfectly consistent to say that Sherlock Holmes is a fictional detective and yet Sherlock Holmes does not exist. Suppose this is in fact true. Then (Gen) above predicts that Sherlock Holmes also *exists among fictional detectives*, since this reduces to saying that Sherlock Holmes is a fictional detective. The domain specifier ‘among fictional detectives,’ can thus shift the domain of quantification to one that includes “new” things that did not

²⁵More generally, I am not assuming that the non-logical constants of the language, were we to add any, are existence-entailing. This idea is known as the *being constraint*. It is usually expressed via the schema $Fx \rightarrow E!x$, where the admissible instances are those where F is a non-logical constant. Notice the claim that *logical* constants express existence-entailing properties is a non-starter if we accept, as I do, a classical propositional logic for Boolean connectives, since under that assumption instances of excluded middle involving a non-existing proposition would be false.

previously exist, so to speak.²⁶ Domain specification is more general than domain restriction, as the latter can be defined from the former: to restrict the domain to just the F s is to specify the *existing* F s as the new domain. Likewise, we can define an operation of domain *expansion*: to expand the domain to include the F s is to specify *being an F or existing* as the new domain.²⁷

(Gen) is the core schema capturing the semantic function of domain specifiers. From it, we can derive further principles describing how domain specifiers interact with quantifiers themselves. If to exist among the F s is simply to be an F , then the result of combining $\text{At}(F)$ with the quantifier $\forall_\sigma$ (where $\sigma \rightarrow t$ is the type of F) should be a quantifier that ranges precisely over the F s. Writing $\forall^F$ to abbreviate $\text{At}(F)(\forall)$, this means we should be able to derive:

$$\forall^F x Fx \quad (5)$$

$$\forall^F x Gx \rightarrow (Fx \rightarrow Gy) \quad (6)$$

Both these principles can indeed be derived from (Gen) together with other auxiliary principles introduced shortly.²⁸

Again, since F need not be existence-entailing, the range of the quantifier $\forall^F$ need not be a *restriction* of that of $\forall$. Quantifiers constructed via domain specifiers can have larger, smaller, or incomparable domains of quantification. In particular, as anticipated earlier, even though $\forall$ itself does not have a classical quantificational logic, we can *define* a quantifier that does by specifying a universal property as the new domain of quantification. By a universal property I do not merely mean a property that everything in the domain has, but rather a property the not having of which is inconsistent. The universal property *being such that* $\top$ (at a given type) can always be defined in purely logical terms, and I will use this property to introduce my classical quantifiers:

$$\forall^\top := \text{At}(\lambda x. \top)(\forall).$$

I will henceforth adopt the convention of writing English quantifiers in boldface, as in **everything**, when I intend the classical reading.

Intuitively, our primitive free quantifiers behave like syntactic restrictions of the defined classical quantifiers just introduced. For example, it could be that although there are no fictional detectives, **there are** fictional detectives, because Sherlock Holmes is a fictional detective. To say that there are no fictional detectives just is to say that **there are** no fictional detectives who *exist*, in the sense of $\text{E}!$. The picture is similar to that put forward by contingentists or actualists who theorize with a pair of quantifiers, one ranging over the “inner” domain of things that in

²⁶More generally, this domain expansion effect can happen with any domain specifier $\text{At}(F)$ where F is not existence-entailing.

²⁷Formally,

$$\text{At}^\cap(F) := \text{At}(\lambda x. Fx \wedge \text{E}!x), \quad \text{At}^\cup(F) := \text{At}(\lambda x. Fx \vee \text{E}!x).$$

²⁸From the same principles we can also derive a schema to the effect that the quantifier $\forall^F$ is *normal*, i.e., closed under modus ponens:

$$\forall^F x (Gx \rightarrow Hx) \rightarrow \forall^F x Gx \rightarrow \forall^F x Hx. \quad (7)$$

These three axioms, together, capture the standard inferential role of a free quantifier that ranges over precisely the F s.

fact exist and one over the “outer domain” of merely possible entities.²⁹ In fact, a quantifier $\forall^F$ always behaves like the syntactic restriction of the defined classical quantifier by the predicate F . In other words, to be in the range of the quantifier $\forall^F$ is simply to **exist** and to be F .

This settles how domain specifiers handle our primitive quantificational idioms. But (Gen) alone does not settle how domain specifiers should interact with arbitrarily complex terms. Consider the property of *being a lover*, in the sense of *loving someone*: $\text{Lover} := \lambda x. \exists y. Lxy$. What is it to be a lover *among the F s*? Intuitively, it is to love **someone** *who is F* : once the domain is reset to the F s, the quantifier ‘someone’ implicit in being a lover ranges over the F s only. (Gen) alone does not deliver this result. It tells us what domain specifiers do to existence predicates, and derivatively to quantifiers, but not what they do to entities, like being a lover, that involve quantification only indirectly. To understand that, we can interpret domain specifiers in terms of *metaphysical substitutions* (Bacon, 2019, 2020).

This notion is easiest to introduce by pretending, for a moment, that reality is structured the way language is. On this pretense, the proposition that Alice loves someone is a complex entity built by applying the quantifier $\exists$ to the property of being loved by Alice, La , which is in turn built by applying the relation L to Alice. If reality were structured this way, there would be a well defined operation of *substituting* one constituent of a proposition (or property, or relation) for another. Substituting Bob for Alice in the proposition that Alice loves someone yields the proposition that Bob loves someone; substituting hate for love yields the proposition that Alice hates someone. A domain specifier, on this picture, is understood as performing a special kind of substitution: $\text{At}(F)$ replaces each quantificational constituent with its counterpart that ranges over the F s—the quantifier $\exists$ with $\exists^F$, the existence predicate $\text{E}!$ with F , and so on in accordance with (Gen)—while fixing every other constituent. The proposition that among the F s, Alice loves someone, then, equals the proposition that Alice loves someone ^{F} (Figure 1). By the same token, to be a lover among the F s is to love someone ^{F} , as desired.




FIGURE 1. Domain specification as substitution: $\text{At}(F)$ replaces the quantifier $\exists$ with $\exists^F$ and fixes every other constituent.

The structured picture is suggestive, but it should not be taken literally. If reality really were structured like language, we would expect that whenever FA is identical to GB , F is identical to G and A to B . But that conflicts with the identification of $\beta\eta$ -equivalents: $\neg\neg p$ is $\beta\eta$ -equivalent to $(\lambda q. \neg\neg q)p$, yet negation is

²⁹There are thorny philosophical questions regarding the interpretation of inner and outer quantifiers having to do with ontological commitment, the substantivity of metaphysical disputes, and the intelligibility of primitive higher-order quantification, both in the contingentist/actualist setting and in the quantificationalist setting. While I cannot satisfactorily address all of them here, I will briefly discuss some in Section 4.

not identical to double negation. Worse, the *Russell–Myhill paradox* shows that the principle is inconsistent in classical higher-order logic. The argument carries over to free logics where a classical quantifier is definable, such as the present one.³⁰

The inconsistency of the quasi-syntactic picture does not imply that the notion of a metaphysical substitution is unintelligible. Instead of defining it in terms of notions like constituency supplied by the quasi-syntactic picture, we can take it as primitive. Its defining feature is that it respects the way that higher-order entities are combined to form new ones: a metaphysical substitution is a kind of transformation on higher-order entities of all types that *commutes with application*, in the sense that performing it on MN yields the same result of first performing it on M and N separately, then applying the results. This idea is intelligible independently of the assumption that reality has quasi-syntactic structure—indeed, it can be made sense of in fairly coarse-grained models of higher-order logic.

I propose, then, to interpret domain specifiers as follows. Each property F determines a metaphysical substitution that “moves” interpretations of existence predicates and quantifiers in accordance with (Gen) and the quantifier principles above, and that fixes everything that has nothing to do with quantification. I will call these substitutions *quantificational substitutions*. How quantificational substitutions act on arbitrary entities is entirely fixed by these facts and the requirement that it commute with application. A term of the form $\text{At}(F)(M)$ expresses the result of executing the quantificational substitution determined by F on M .

This idea of an entity that has nothing to do with quantification is important and will come up repeatedly. I call such entities *stable*. We have seen many examples of properties that are not stable. Some involve quantification directly: existence properties and quantifiers themselves. Others do so less directly, like the property of loving someone. Others still indirectly involve quantification even though we do not express them using overtly quantificational language. For example, the term ‘to eradicate’ does not contain overt occurrences of quantifiers, but its meaning involves quantification: to *eradicate* a disease is to eliminate *all* (or perhaps *nearly all*) instances of it. Stable entities are supposed to be entities that do not involve quantification even in such indirect ways, and should therefore be fixed by all domain specifiers: whenever M is stable, the identity $\text{At}(F)(M) \equiv M$ always holds.

I consider *combinators* and Boolean operations to be paradigmatic examples of stable entities. A combinator is the denotation of a closed term that contains no constants, for example $\lambda p.p$ (truth) or $\lambda X^{\sigma \rightarrow \tau} Y^{\tau \rightarrow \rho} z^{\sigma}.Y(Xz)$. We can describe the behavior of these entities without appealing to quantification at all. To compose

³⁰This argument is originally due to Russell (1903, Appendix B) and Myhill (1958) and recently revived, e.g., in Uzquiano (2015, 2023); Dorr (2016); Goodman (2017); Fritz et al. (2021); Bacon (2023b). Here is a sketch. Given an arbitrary proposition p , consider the property of being *heteropredicative* (with respect to p): this is a property that a proposition q has when every property it predicates of p is a property that q itself lacks:

$$\text{Het} := \lambda q.\forall X(q \equiv Xp \rightarrow \neg Xq).$$

We can prove in classical higher-order logic that the proposition $\text{Het}p$ is not heteropredicative: there is some property of propositions Y that $\text{Het}p$ predicates of p that $\text{Het}p$ itself has. This is a counter-example to the structural principle mentioned at the beginning of the paragraph: $\text{Het}p \equiv Yp$, yet $Y(\text{Het}p)$ and $\neg \text{Het}(\text{Het}p)$, showing Het and Y must be different properties.

X with Y is just to apply Y after X ; the conjunction of P and Q is true precisely when P is true and Q is true.³¹

I will also assume that the higher-order identity relations $\equiv_\sigma$ are stable. This might feel more contentious. In higher-order metaphysics, it is common to define identity as the relation two entities bear when they share *all* their higher-order properties. In a free logical setting, we need to be careful with a definition of this sort: since universal generalizations fail to imply all their instances, it might be that two entities share all their higher-order properties and yet one is F while the other one is not, because F does not exist. This problem, of course, can be circumvented in our system by replacing the defined classical quantifier for the free one: identity as the relation two entities bear when they share **all** their higher-order properties. But the defined classical quantifiers are provably stable, owing to the fact that they are defined by specifying a stable property, $\lambda x. \top$, as the domain of quantification within them. I will return to this point towards the end of the section.

Whether any other entities are stable is a question that the framework leaves open. An attractive idea is that individuals (entities of type e) should always be stable; they simply are not the sort of thing that can involve quantification. Another interesting thought is that *fundamental* entities should be stable. This is especially fitting in the case of propositions if we think of fundamentality in terms of *grounding*. Grounding theorists typically regard universal and existential generalizations as grounded in their instances and therefore non-fundamental, so at least those kinds of non-stable propositions cannot be fundamental. Finally, I am attracted to the view that stable entities are closed under a *comprehension* principle: roughly, **every** proposition, property or relation is (classically) coextensive with **some** stable proposition, property or relation. I will not try to settle these matters here.

So much for stability. Once we embrace the substitutional picture, a number of further principles about domain specifiers naturally flow from it. Of course, domain specifiers should commute with application, since quantificational substitutions do:

$$\text{At}(F)(MN) \equiv \text{At}(F)(M)\text{At}(F)(N). \quad (\text{App})$$

In quantificationalist terms: to specify the domain within MN is the same as to specify it within M and within N , then apply the former to the latter. Application contributes no quantification of its own, so whatever quantification MN involves must come from M or N .

Next, we have two principles describing *closure conditions* on substitutions. Substitutions, it is natural to think, form a rich family. First, there is always an *identity* (or trivial) substitution, which moves nothing. This idea has a natural quantificationalist interpretation. If the current domain of quantification is given by *the existing things*, there should be a quantificational substitution that specifies *the*

³¹You might think that Boolean operations actually do involve quantification indirectly. For example, you might think that negation is the property of materially implying *everything*, or that disjunction is the relation two propositions bear when *one of them* is true. If so, you will think Boolean operators are not stable, but you will also think that their logic is non-classical. For example, if neither P nor its negation exist, then $P \vee \neg P$ will be false. These glosses on Boolean operations are extensionally adequate if we replace the quantificational phrases above with their boldface classical counterpart, and intensionally adequate if we assume that our defined classical quantifiers satisfy domain-invariabilism. As I note later, the latter claim is a theorem of my system, a consequence of the claim that Boolean operations, and therefore $\top$, are stable.

existing things as the domain of quantification within its argument. That should not change the domain of quantification at all: to be F among the existing things just is to be F . More generally, we should have the schema

$$\text{At}(\text{E!})(M) \equiv M, \quad (\text{Id})$$

where M can have any type.

Second, given any two substitutions, there is always another substitution which *composes* the two, corresponding to applying one after the other. This idea, too, generalizes to quantificational substitutions. Every quantificational substitution is determined by a property. Thus, if we say that quantificational substitutions should be closed under composition, we must be able to find a property that determines the composition of the quantificational substitution determined by F with the one determined by G . The property *being G among the F s* is the most natural choice. In other words, to specify the F s and then the G s as the domain one after the other is the same as to specify *the G s among the F s* as the domain in one step. Formally,

$$\text{At}(F)(\text{At}(G)(M)) \equiv \text{At}(\text{At}(F)(G))(M). \quad (\text{Comp})$$

This principle helps us resolve a number of questions concerning nested domain specifiers. Here is one example. Naively, you might have thought that specifying the same domain “twice over” is redundant: something is true among the F s just in case among the F s, it is true among the F s. (Comp) explains why this is not so: whenever F is not stable this redundancy idea can fail, because to be F need not be the same as to be F among the F s. To illustrate, consider again the property of being a lover (*loving someone*). Suppose Alice loves Bob and Bob loves Carol, but Carol does not love anyone. Then Alice and Bob are both lovers, yet only Alice is a *lover among lovers*. So, it is true that Bob exists among lovers, but false that, among lovers, he exists among lovers.

Another consequence of (Comp) is that whatever you get by specifying a stable property as the domain of quantification must itself be stable. In particular, as I anticipated earlier, specifying the domain to the F s within $\text{At}(\lambda x. \top)(\forall)$ amounts to specifying $\text{At}(F)(\lambda x. \top)$, which is just $\lambda x. \top$, as the domain within $\forall$. The stability of defined classical quantifiers is an important feature of the framework. Domain specifiers do not affect all quantifiers equally. There is one collection of quantifiers, our primitive free quantifiers, whose theoretical role is, in part, to be maximally *flexible*: their domain can be specified arbitrarily. When I informally say that an entity “involves quantification,” what I mean is that it involves the kind of quantification that the primitive free quantifiers are doing. Other quantifiers, be they definable from the primitive free quantifiers or not, may have different theoretical roles. Classical quantifiers, in particular, are intended to be maximally *general*, rather than maximally flexible, and remain so in every context. Their stability secures this: because the defined classical quantifiers are stable, they satisfy classical quantificational logic within the scope of arbitrary domain specifiers.

With these principles in place, we can derive the original verdict that to love someone among the F s just is to love **someone** who is F if we also assume that the *loving* relation is stable. Likewise, we can derive that the proposition that, among the F s, Alice loves someone is the same as the proposition that, Alice loves **someone** who is F , given the additional assumption that Alice, too, is stable. Without these stability assumptions, matters are more complicated, though that is

to be expected. What results we get depends on how, exactly, we take the *loving* relation and Alice herself to involve quantification.

I close this section by returning to the syncategorematic status of domain specifiers. The main philosophical reason why domain specifiers are syncategorematic is that they describe a way of combining higher-order entities that irreducibly differs from application. If domain specifiers were categorematic terms, then to specify the domain of quantification would amount to *applying* a higher-order entity *to* another. But that is not how we should think of it. Rather, to specify the domain is to *substitute* a higher-order entity *for* another. Metaphysical substitutions are transformations that *preserve* applicative structure, not entities that can themselves be applied to others.

On the structured pretense, the difference between applying and substituting is easy to see. The proposition that Alice runs has two constituents: Alice and the *running* relation. Applying negation to the proposition that Alice runs yields a proposition with three constituents: the original two, plus negation. By contrast, substituting Bob for Alice in that proposition yields a proposition with exactly as many constituents, in which neither Alice nor anything playing the role of a “substituter” occurs. This is not to say that substitution never increases constituent count: if we substitute *walking fast* for *running* in the proposition that Alice runs, the result has three constituents, with the extra one being the property of relations expressed by the adverb ‘fast.’ But application *always* preserves the original constituents, and increases constituent count whenever the entity that is being applied has constituents itself. Substitution, by contrast, need not preserve the original constituents and does not increase constituent count when the substituted entity and the entity being substituted for have the same number of constituents.

I have urged, however, that the structured pretense not be taken literally, so this argument, while suggestive, is certainly not decisive. It is not entirely useless either: quasi-syntactic accounts of metaphysical structure have been shown to be consistent when appropriately restricted, for example to *fundamental* entities only (Bacon, 2023b). The argument above is effective for quantificationalists who do buy such restricted accounts.

There is a stronger, if less suggestive, argument to the effect that if we think of domain specifier composition in accordance with (Comp), like the substitutional picture suggests we do, then we should not think that domain specification is a matter of applying a higher-order entity to another. Here is the rough idea. Suppose that specifying the *F*s as the domain within *M* simply amounted to applying some entity s_F to *M*. Then s_F is an entity like any other, and further domain specifiers must act on it too. Either s_F is stable or it is not. If it is, then by (App) specifying the *G*s as the domain within $s_F M$ yields s_F applied to $\text{At}(G)(M)$: the two specifications would *commute*, so that resetting the domain to the lovers and then to the unloved would be the same as doing so in the reverse order. But they plainly differ: the first should give the domain of those who are unloved among the lovers, the second that of those who love someone unloved. If s_F is not stable, then specifying the *G*s within $s_F M$ acts on *M* *twice over*: once directly, and once again through its effect on s_F , which is itself sensitive to the domain. Neither option gives the verdict of (Comp), on which the outer specifier acts on *M* only once, by acting on the inner domain.

This argument gives us strong reason to deny that quantificational substitutions combine with their arguments through application, and in turn to embrace the syncategorematic treatment of domain specifiers. In Section 6.1 I will present a more formal version of this argument and show it actually pushes further than that. The syncategorematic treatment of domain specifiers only does justice to the substitutional picture in the presence of non-standard restrictions on λ -abstraction. But we can bracket these issues for the time being.

4. OUTLINE OF A RESEARCH PROGRAM

I do not have a knock-down argument for quantificationalism, and I do not really know what such an argument would even look like. As I see it, the case for quantificationalism must be abductive: it should be accepted, if at all, because it provides a fruitful framework for theorizing about a range of philosophical topics, provided the benefits outweigh the costs. What I want to do in this section and the next is to try and convince the reader that there is a research program with legs here, which once carried out can serve as the basis for an abductive argument in favor of quantificationalism.

In this section I sketch a few possible applications of quantificationalism. My goal here is to convey the breadth of the research program, so I will stick to programmatic overviews and leave details for future work. In the next section I will present an additional application of quantificationalism in more depth, as a proof of concept: a quantificationalist approach to the notion of *predicativity*.

4.1. Ontological Disputes. A first application is to *ontological disputes*, namely disputes about what exists. The paradigm is the dispute between the *universalist*, who holds that any things whatsoever compose something, and the *nihilist*, who holds that nothing is composed of anything, so that there are no tables, only simples arranged tablewise (Lewis, 1986; van Inwagen, 1990; Hawthorne and Cortens, 1995).

Witnessing their prolonged intractability, some philosophers have thought we should be *deflationists* about many ontological disputes. There is no real disagreement between the universalist and the nihilist. Their dispute is *merely verbal*, much like the dispute about whether a hot dog is a taco. Others, the *realists*, have resisted this conclusion. The disagreement between universalists and nihilists is substantive. The fact that we have not found out which side is right just shows that the question is hard, not that it has no right answer. Quantificationalism has something to offer to both the deflationist and the realist, though I think that, perhaps surprisingly, the most interesting upshots are for the realists.

The best known version of deflationism about ontology is *quantifier variance* (Hirsch, 2002, 2010). The universalist and the nihilist are talking past each other because they mean different things by the quantifier ‘there is.’ Each of them speaks truly in their own language. You might think quantificationalism is an especially useful framework for quantifier variantists, since it makes generating alternative quantifier meanings very easy: for any property F , $\text{At}(F)(\forall)$ is a quantifier ranging over precisely the F s. There is something to this thought. Sider (2009, §5) argues that the quantifier variantist owes us translations between the universalist’s language and the nihilist’s that preserve logical form, mapping quantifiers to quantifiers. This is easily done in one direction: the universalist simply interprets the nihilist’s quantifier as a restriction of their own to the simples. The other direction, Sider claims, is harder, since standard logical frameworks do not provide

the tools for explicitly defining quantifier *expansions*. The best the nihilist can do is paraphrase universalist quantification in terms of plural quantification over simples, but this at best yields a truth-preserving sentence-to-sentence mapping, not one from the universalist’s quantifier to another quantifier. Quantificationalism supplies exactly the missing tools. The nihilist can define the universalist quantifier by specifying the property *existing or being composite* as the new domain of quantification within their own primitive quantifier. According to the new quantifiers, there are tables, even though in fact there are none.

That being said, I think the most interesting way that quantificationalism re-frames the debate over substantivity is on the realist’s side. Realists of a certain stripe argue that we can engage in substantive ontological disputes because, even though different quantifier meanings verify and falsify different ontological theses, we can single out a *special* quantifier meaning and decide to conduct ontology in terms of it. Ontology deals with what exists in the sense of the special quantifier.

The realist should say something about what makes the special quantifier special. A particularly attractive idea is that the special quantifier can be characterized purely by its inferential role: the special quantifier is just the *classical* quantifier (Williamson, 1988, 2003; Dorr, 2014; Bacon, 2023a) [CHECK DORR].³² We can speak of “the” classical quantifier because any two expressions that satisfy the inferential role of a classical quantifier must be logically equivalent (Harris, 1982). Insofar as two sides of an ontological dispute in fact are using their quantifiers classically, they must be employing the same notion of existence up to logical equivalence.

A worry about this approach is that many philosophers do not use their quantifiers this way when doing ontology. We already saw in Section 3.2 that for contingentists and actualists in the metaphysics of modality, the kind of quantification that matters for ontology is not classical but merely free. At the same time, these philosophers do employ classical quantifiers in their theorizing without viewing it as ontologically committing, suggesting they are not simply mistaken about what the correct logic of their notion of existence is. If so, then the dispute between contingentists and necessitists over, say, whether Wittgenstein’s merely possible first child exists would seem to count as merely verbal.

That might be a welcome result for some, but there might also be realists who want to resist this conclusion. The inferential role standardly associated with free quantifiers does not uniquely pin down a single meaning up to logical equivalence, so adapting the inferentialist strategy just sketched to ensure consistency with the idea that the special quantifier is free is not straightforward. Quantificationalism provides the resources to do just that, in a novel way. While the *standard* inferential role associated with free quantifiers does not uniquely pin down their meaning, the richer inferential role that quantificationalists assign to their primitive free quantifiers does.³³ The quantificationalist can say that the special quantifier is distinguished in virtue of interacting with domain specifiers the right way.

³²A perhaps better known approach is that of Sider (2009, 2011), who suggests instead that the special quantifier is the *fundamental* one, the one that “carves reality at its joints.” The worry I am about to raise does not apply to this approach.

³³Williamson (1988) has a similar approach. He suggests contingentists should accept the *being constraint*, a schema to the effect that each primitive predicate of the language expresses an existence-entailing property. This, of course, makes the meaning of existence unique only relative to a choice of signature.

The relevant uniqueness results come from two simple ideas about domain specification. First, that *to exist among the F s* is just to be an F (**Gen**). And second, that there is a trivial way of specifying the domain of quantification, namely by specifying *the existing things* (**Id**). Whenever ‘exists’ is used this way, a single meaning is pinned down up to logical equivalence.

4.2. Quantificationalist Modal Realism. A second application concerns the reduction of modality to quantification.

Lewis (1968, 1986) famously proposes to reduce modality to quantification over possible worlds: possibly, P just in case P is true at w for some possible world w . Lewis suggests that the modifier ‘at w ’ works

mainly by restricting the domains of quantifiers in its scope, in much the same way that the restricting modifier ‘in Australia’ does. In Australia, all swans are black—all swans are indeed black, if we ignore everything not in Australia. [...] At some strange world $[w]$, all swans are blue—all swans are indeed blue, if we ignore everything not part of w . (Lewis, 1986, p. 5)

Lewis does not ultimately *analyze* the modal operator ‘possibly’ as a combination of a quantifier over worlds and a domain restricting phrase. He instead proposes that modal operators and phrases of the form ‘at w ’ themselves be *paraphrased away* using quantification over worlds and syntactic restrictions on quantifiers.³⁴ To illustrate, Lewis paraphrases (8) as (9)

(8) Possibly, all swans are black.

(9) There is a possible world w such that every part of w that is a swan is black.

In the paraphrase, there is no operator expression ‘at w ’ that functions as a quantifier domain restrictor. The idea that modal operators work partly as quantifier domain restrictors is instead implemented by translating sentences that occur within the scope of a modal operator in the target language by syntactically restricting quantifiers to parts of a specific world. In Sider’s (2009) terminology from the previous subsection, Lewis’s translation does not preserve logical form.

In quantificationalism, however, Lewis’s informal suggestion that modal operators are the composition of a quantifier restricting phrase with a quantifier over worlds can be taken literally. We can articulate a quantificationalist version of modal realism where possibility is equivalent to classical existential quantification over possible worlds followed by a domain specifier that restricts the domain to parts of the witnessing world:

$$\Diamond P \leftrightarrow \exists^\top w(\text{World } w \wedge \text{At}(\lambda x.x \preceq w)(P)).$$

Here the phrase ‘at w ’ is rendered as the domain specifier $\text{At}(\lambda x.x \preceq w)$, where $\lambda x.x \preceq w$ is the property of being a part of w . This expression is not a notational device for a paraphrase schema. Though syncategorematic, it is part of the language and really occurs in operator position. The equivalence above can be strengthened to a domain-invariably strict equivalence or even a higher-order identity.

³⁴This is the proposal that emerges from Lewis (1968), though Lewis (1986) is skeptical that the paraphrase project can succeed and proposes instead to theorize directly in terms of modal realist ideology rather than treating it as a device to regiment ordinary modalizing.

The quantificationalist approach allows us to make literal sense of other remarks of Lewis that, arguably, do not survive the counterpart-theoretic paraphrase implementation. Lewis explicitly suggests that the phrase ‘at w ’ modifies not only the interpretation of quantifiers, but also those of phrases not overtly quantificational but which “reveal quantification under analysis,” such as verbs like ‘to invent’ or my previous example ‘to eradicate,’ and even of proper names (Lewis, 1986, p. 6). But the counterpart-theoretic translation of Lewis (1968) fixes all primitive predicates and only applies to sentences of *pure* quantified modal logic, which lacks individual constants entirely. The only way of making sense of these suggestions in counterpart theory is to say that one should first analyze ordinary language to make implicit quantification syntactically explicit and translate proper names along descriptivist lines, then apply the counterpart-theoretic paraphrase to the result (Divers and Parry, 2018). But in quantificationalism domain specifiers can directly combine with predicates and names and make a non-trivial semantic contribution to the result even if said predicates and names contain no overt occurrences of quantifiers.

Besides taking these suggestive Lewisian remarks at face value, quantificationalist modal realism can also be expected to fare well with respect to the problem of *advanced modalizing* that plagues the Lewisian paraphrase approach. This consists in a collection of puzzles arising when the Lewisian paraphrase is applied to sentences of a language containing both ordinary modal operators and ‘advanced’ modal vocabulary such as ‘possible world,’ ‘counterpart,’ and so on. A simple example: Lewis is committed to the claim that there are at least two worlds, so he should also accept that possibly, there are at least two worlds. But the Lewisian translation of the latter claim is that some world contains two worlds as parts, which Lewis rejects.

This problem arises because Lewis thinks of quantifier domain restriction as a syntactic matter. The quantifier in ‘there are at least two possible worlds’ is intended to have an unrestricted reading, but translating the modal operator forces a syntactic restriction. In quantificationalism, on the other hand, we understand domain specification semantically. Unrestricted quantification over worlds is rendered via the defined classical quantifier $\exists^\top$, whose domain cannot be re-specified because $\lambda x. \top$ is a stable property. Thus, the claim that possibly, there are at least two worlds is rendered as

$$\Diamond \exists^\top wv (\text{World } w \wedge \text{World } v \wedge w \neq v).$$

In the presence of suitable stability assumptions this turns out to be equivalent to the non-modal claim that there are at least two worlds:

$$\exists^\top wv (\text{World } w \wedge \text{World } v \wedge w \neq v).$$

The quantificationalist solution to this puzzle resembles the *redundancy* solution offered by Divers (1999, 2014): the Lewisian paraphrase should only be applied to sentences where “quantification is world-restricted.” Otherwise, modal operators should be treated as *redundant*. But on the quantificationalist approach, redundancy is *derived* from independently motivated principles in the logic of domain specifiers, rather than stipulated.

4.3. Further Topics. Finally, I briefly discuss two additional areas where quantificationalism might apply, though I should emphasize that its prospects in these domains are more speculative. The first is the semantics of *exceptive constructions*

in natural language. I introduced domain specifiers by analogy with phrases like ‘with the exception of the United States.’ It is natural to wonder whether this analogy could rise to the level of an *analysis*.³⁵ Linguists distinguish *connected* exceptives, such as ‘every student but Kim,’ from *free* exceptives, such as ‘except for Kim.’ There is some evidence that free exceptives are syntactically sentence adjuncts (Hoeksema, 1987; Von Stechow, 1993), which could perhaps be used to motivate an analysis that treats free exceptives as behaving semantically like sentential operators. Von Stechow (1993) insists that free exceptives cannot be interpreted “at the sentence level” because that would predict that a sentence like ‘Except for John, everybody likes John’ is uninterpretable, which it is not. To a quantificationalist who interprets free exceptives as domain specifiers, however, this sentence is perfectly interpretable: we just need to say that *liking* is not existence-entailing in its second argument place.

The second area, more promissory still, is *potentialism* in the philosophy of mathematics. Potentialists hold that the universe of sets is never complete: any domain of sets can be expanded to a larger one (Fine, 2006b; Linnebo, 2013, 2018). They typically articulate this idea using modal operators, which raises the question of what sort of modality is at issue. Quantificationalism suggests an answer: the potentialist’s modality might just be identified with a restriction of domain-occasional truth to domains that expand the current one only by adding “new” mathematical objects. Developing this idea would presumably require the generalization of the framework sketched in Section 3, on which not every property determines a domain. Otherwise, the classical quantifier $\exists^\top$ would provide a domain that cannot be expanded, which is exactly what the potentialist denies.

5. METAPHYSICAL PREDICATIVITY

There is another, to my mind more ambitious application of quantificationalism that I want to discuss. Quantificationalism offers a new perspective on the notion of *predicativity*. I discuss this application of quantificationalism in greater detail than the previous ones. If discussion of the previous applications was aimed at demonstrating the breadth of the quantificationalist research program, my aim here is to give a proof of concept of what a concrete application of quantificationalism might look like.

I will propose a particular notion of predicativity, *metaphysical predicativity*, on behalf of the quantificationalist. It is a *worldly* notion: propositions, properties and relations, as opposed to linguistic entities like definitions or instances of axiom schemata, are said to be metaphysically predicative or impredicative. Metaphysical predicativity earns its keep by helping articulate novel solutions to a number of paradoxes—or, more carefully, solutions that have been gestured at for a long time but arguably never precisely articulated.

³⁵I want to stress again that neither the intelligibility of domain specifiers nor the truth of quantificationalism rests on this. The meaning of domain specifiers is not given by that of any natural language expression: it is fixed by their theoretical role, namely the logic described in Section 3 and its interpretation in terms of quantificational substitutions. This is why I insisted that pronouncing $\text{At}(F)$ as ‘among the F s’ is merely a convention. If exceptives turned out to be best analyzed as, say, modifiers of determiners rather than sentence adjuncts, that would not decide the question whether quantificationalism is true, just as the question whether temporalism is true would not be decided by the discovery that the English expressions ‘always’ and ‘sometimes’ are not sentential operators (Sections 1 and 2).

5.1. Impredicative Definitions. A definition is called *impredicative* when it defines an entity by means of quantification over a domain that includes the very entity being defined. Otherwise, it is called *predicative*. For concreteness, consider:

- (10) The intersection of A and B is the set containing all and only those elements that A and B have in common.
- (11) To be a natural number is to possess every property had by 0 and closed under the successor function.
- (12) The *Russell set* is the set containing all and only the sets that do not contain themselves.
- (13) The *truth-teller* is the proposition that there is a true proposition.

Intuitively, (10) is a predicative definition. It defines the intersection of A and B by quantifying over the common elements of A and B , but the intersection of A and B itself is not one of these common elements.

On the other hand, (11) is impredicative. It defines the property of being a natural number by quantifying over all properties had by 0 and closed under the successor function. But being a natural number is itself one such property. (12) is also impredicative. It defines r by quantifying over all sets that do not contain themselves. But r itself is (and, paradoxically, is not) such a set. (13) is also impredicative, as it defines a proposition by quantifying over all true propositions, including itself.

The notion of predicativity originates in [Russell \(1903\)](#), influenced by the work of Poincaré. Russell is led to the notion of predicativity by the idea that the class-theoretic paradoxes arise because of viciously circular quantification, and that the proper solution to the paradoxes is to prohibit such quantification. This prohibition, now known as the *vicious circle principle* (VCP), is articulated by Russell in different ways. Two explicitly mention definitions.

- (VC1) If, provided a certain collection had a total, it would have members only definable in terms of that total, then the said collection has no total. ([Whitehead and Russell, 1910](#), p. 20)
- (VC2) No totality can contain members defined in terms of itself. ([Russell, 1908](#), p. 237)

These principles, while far from transparent, have long been interpreted by logicians as asserting a ban on impredicative definitions. On this interpretation, for example, the argument underlying Russell's paradox is fallacious because the impredicative definition of the Russell set is illegitimate.

The distinction between predicative and impredicative definitions, as just now glossed, feels intuitive enough. But making it precise is not straightforward. Different authors have done so in different ways, and there is no single well-established notion of an impredicative definition in the literature.³⁶ Here I put forward one plausible precisification of the notion of impredicative definition, without making any claim that it is the only possible or the best one for all purposes.

³⁶[Feferman \(2005\)](#) is the canonical survey. In fact, the very notion of an impredicative definition does not really appear in the theories predicativists build. The ramified theory of types of [Russell \(1908\)](#); [Whitehead and Russell \(1910\)](#) can be seen as a logical framework that makes impredicative definitions impossible, and so does not have the expressive resources to itself distinguish between predicative and impredicative definitions. Others (e.g., [Weyl, 1994](#)) implement predicativity as a way of restricting acceptable instances of comprehension schemata.

My approach is to take the intuitive gloss of an impredicative definition, that of a definition that quantifies over a domain that includes the entity being defined, as the starting point, and to make precise sense of what it means for a linguistic expression to quantify over a domain that includes a given entity. This requires some care. Naively, we might propose to identify the domains over which a definition quantifies with the values that the variables bound by the quantifiers occurring within it can take. That is, a definition quantifies over a domain D (that includes a) iff it contains a quantifier binding variables that can take any value from D (and D contains a). Call this the *naive proposal*.

The naive proposal is too crude. The main problem is that in a language with unrestricted quantifiers whose logic is classical, any definition of an entity of type σ that contains quantifiers binding variables of type σ counts as impredicative. But the distinction between predicative and impredicative definitions makes intuitive sense in such languages. For concreteness, consider (10). It is intuitively a predicative definition of the intersection of A and B . Yet when regimented in a first-order language for set theory, it will contain an occurrence of the unrestricted first-order quantifier ‘all.’ In the usual interpretation of set theory, variables bound by this quantifier can take any set as their value, including the intersection of A and B . On the naive proposal, (10) is just as impredicative as (12).

Building on a suggestion of Chihara (1973, p. 6), I suggest we revise the naive proposal by considering not just what the quantifiers in fact range over, but also at how what a definition picks out changes as we vary what the quantifiers range over.³⁷ Impredicative definitions quantify over a domain that includes the entity being defined in the sense that they are *quantificationally sensitive* to their definiendum: *what* they define can change depending on whether we add or remove the definiendum from the domain of quantification.

To illustrate, consider (11). Suppose we force the second-order quantifiers in the definiens to range only over the domain of properties that apply to finite initial segments of the natural numbers. This domain does not include any hereditary property, so (11) defines a property that vacuously applies to everything. Once we add the property of being a natural number back to the domain, on the other hand, the definiens picks out a property that does not apply to everything.

Or consider (12). Russell’s paradox arises from the crucial assumption that the Russell set itself falls within the range of the quantifiers occurring in its definition.

³⁷Chihara discusses the proof of Cantor’s theorem, which shows that there is no surjection f from a set A to its powerset $\wp(A)$. The proof is by contradiction. We assume there is such a surjection f , which in set theory is represented by a set of ordered pairs. We then define a set B via the following membership conditions:

$$\forall x(x \in B \leftrightarrow x \in A \wedge \forall y((x, y) \in f \rightarrow x \notin y))$$

Chihara observes:

Let us suppose that in our specification of B , ‘ x ’ and ‘ y ’ are restricted to relatively small totalities, viz. A and $\wp(A)$ respectively (clearly, we cannot restrict these variables to smaller ranges without specifying a different set from the one intended). Even then, the specification would be impredicative since B would be a member of the range of ‘ y ’.

The crucial observation here is that there is a domain, namely the union of A with $\wp(A)$, such that if we let the quantifiers in our definition range over that domain, the definition picks out its intended definiendum. But if we remove the intended definiendum from that domain, the definition ends up picking out a different set.

Without this assumption, the argument cannot be run. And clearly, there are lots of choices of domains not including the Russell set, such that if we forced the quantifiers in the definiens to range over these domains, we would get something other than the Russell set as the result of the definition.

Analogously for (13). If we force the quantifiers in the definiens to range over, say, the actual truths of biology, then (13) defines a proposition that is false at a world where the truths of biology are wholly disjoint from the actual truths of biology. Once we add the proposition that there is a truth back into the domain, however, the definiens picks out a proposition that is true at such a world.

Paradigmatically predicative definitions, by contrast, do not behave this way. While (10) does define different sets depending on what we take the quantifiers to range over, adding or removing the intersection of A and B itself from the domain of quantification cannot change what (10) defines, no matter what we take the domain to be. Sets are individuated extensionally, and the membership conditions specified by (10) are equivalent to those we get by adding or removing the intersection of A with B from the domain of quantification.

5.2. Beyond Syntax. The notion of predicativity I have been discussing is a property of linguistic entities, namely of definitions.³⁸ However, predicativity was not originally meant to be a purely linguistic notion. Following Russell, predicativists hold that some (or, sometimes, all) impredicative definitions are illegitimate because they exhibit a kind of vicious circularity. But this circularity is not purely a matter of linguistic form. Rather, it is the very entity being defined by an impredicative definition—or the class over which the impredicative definition quantifies—that, if it existed, would be involved in a vicious circle of some sort. There is, in other words, a *metaphysical* notion of predicativity underwriting the linguistic one. And we would be better served by theorizing directly in terms of the metaphysical notion instead of taking a detour through its linguistic approximation. This section elaborates on these points.

Why are impredicative definitions illegitimate? (VC1) and (VC2) appear to make ontological claims: there is nothing corresponding to an impredicative definition (VC1), or there is nothing that can only be defined impredicatively (VC2). An impredicative definition, then, is illegitimate because it either always fails to pick out anything, or it is at risk of doing so—because it might be attempting to define something that can only be defined impredicatively.

Why is that the case? If we think that at least *some* impredicative definitions are non-empty, the problem does not seem to lie in the grammar of impredicative definitions themselves. Rather, predicativists suggest that there would be something problematic in the *metaphysics* of the entities that only admit impredicative definitions, if there were any, or in the classes quantification over which can be used to state impredicative definitions. The definition of being a natural number as having every hereditary property is viciously circular, because it defines a property by quantifying over a domain that includes that very property. But if that is the *only* definition of being a natural number, that reveals a way in which

³⁸Philosophers working with the Aristotelian concept of a *real definition* (e.g., Fine, 1994) might think there is a sense in which a definition is not merely a linguistic entity but rather some kind of worldly entity that specifies the nature of something. How predicativity interacts with the notion of a real definition is an intriguing question, but one I will not pursue here.

the property of being a natural number *itself*—or perhaps the class of hereditary properties—exhibits a kind of vicious circularity.

Something like this picture is suggested more strongly by other formulations of the VCP proposed by Russell, which do not mention definability at all. In particular,

(VC3) Whatever in any way *concerns all or any or some* [...] of a class must not be itself one of the members of a class (Russell, 1973, p. 209).

In (VC3), the relation *concerning all or any or some* is borne by an object to a class. No explicit reference to language is made. For brevity, let us call this relation *quantificationally concerning*. The VCP, as articulated by this statement, asserts that this relation cannot enter (presumably vicious) cycles of a particular form: no entity can quantificationally concern a class (set with a total) it is a member of.

In sum, impredicative definitions are illegitimate because, if I can only define an entity by means of quantification over a domain that includes that very entity, this is evidence that the entity itself quantificationally concerns a class that contains that very entity, and *that* is what is metaphysically problematic. The VCP is ultimately a metaphysical thesis, imposing the absence of such vicious circles of quantificational concern as a necessary condition on existence.³⁹

I am suggesting that we can view predicative definability as an attempt to linguistically characterize the worldly property of not quantificationally concerning a class that contains oneself. I will now argue that a predicativist who thinks everything has this property should not regard predicative definability as a good approximation of it. For what counts as predicatively definable is language-dependent in a way that existence is not.

For example, whether the number 1 is predicatively definable depends on whether we work in a signature equipped with a constant symbol for 0 or not. If we are, we can define 1 predicatively as the successor of 0. If we are not, it is most natural to define 1 as the smallest natural number greater than 0, or as the smallest natural number with a predecessor, both of which are arguably impredicative definitions. Yet either 1 exists or it does not, regardless of the language we use to talk about it. More generally, predicative definability is best regarded as a relation rather than a unary property: an entity is predicatively definable *relative to* a language (or perhaps *given* some other entities the existence of which we assume to be unproblematic, see Feferman and Hellman, 2021).

We can introduce unary versions of predicative definability following familiar blueprints. We could say that an entity is predicatively definable *simpliciter* if it is predicatively definable in *some* language. This would overgenerate: anything is predicatively definable in a language equipped with a constant symbol standing for it. Nor would it do to say that something is predicatively definable *simpliciter* if it is predicatively definable in *every* language: the existence of expressively impoverished languages would make this notion far too demanding. For the predicativist who wants to use predicative definability as a necessary condition on existence, perhaps the most promising approach is to say that something is predicatively

³⁹Gödel (1944, p. 135), commenting on (VC3), distinguishes between a description of an entity involving all of a class from the entity itself involving all of a class, supporting the present reading. Skolem (1967, p. 297) writes that what seems puzzling about impredicative definitions of sets is that they presuppose that there are “sets *whose existence* depends upon all sets” (emphasis added).

definable *simpliciter* if it is predicatively definable in some metaphysically distinguished language. That, of course, invites a host of thorny questions about what makes a language metaphysically distinguished in the target sense.

A cleaner approach, if it can be made to work, is to avoid the detour through language altogether. The sort of predicativist I am envisaging wants to assert that no entity quantificationally concerns a class that contains that very entity. The thesis that everything must be predicatively definable feels like a clearer linguistic proxy for this claim, but ultimately falls short of capturing the target thought. I propose, instead, to bypass the notion of predicative definability entirely, and make the thought precise by giving a purely non-linguistic account of what it means for an entity to quantificationally concern a class that contains that very entity. This account consists in a purely non-linguistic generalization of the notion of quantificational sensitivity introduced in Section 5.1.

It should be emphasized that, if we do not assume quantificationalism, this is far from a trivial task. In the linguistic case, the notion of sensitivity is characterized in terms of how manipulating the ranges of the quantifiers occurring in the definition can change what the latter defines. But the notions of a quantifier, of an occurrence of a quantifier, and of the range of a quantifier are all linguistic notions that do not have straightforward worldly analogs.

To be sure, there are substantive metaphysical theories that predict worldly analogs of these notions do exist. The obvious candidate is a quasi-syntactic picture of metaphysical structure, which we already encountered in Section 3.3. On the quasi-syntactic picture, we might think that the proposition that *there is a truth* has a metaphysical structure corresponding to the existential syntactic structure of the sentence ‘There is a truth:’ it is a whole made up of two constituents, a *quantificational constituent* (the higher-order entity expressed by the existential quantifier) and the property of *being true*. A picture of this sort would allow us to de-syntactify the definition of sensitivity by replacing syntactic notions with their quasi-syntactic metaphysical counterparts. Manipulations in the range of a quantificational constituent can be made sense of in terms of substitution of one quantificational constituent for another, like we did when sketching the interpretation of domain specifiers against the backdrop of the quasi-syntactic picture.

As we already know, however, the quasi-syntactic picture is inconsistent. We thus have reasons to explore approaches to de-syntactifying the notion of quantificational sensitivity that do not rely on the assumption that reality has quasi-syntactic structure. Quantificationalism has the resources to do just that.

Roughly, we define an entity M to be *quantificationally sensitive* to another entity N when, at some domain, M is not equivalent to the result of adding or removing N from the domain of quantification. Here, ‘equivalent’ is a placeholder. Any equivalence relation, from coextensiveness to identity, can be plugged in, yielding different notions of quantificational sensitivity.

The operation of adding or removing something from the domain of quantification is made sense of through domain specifiers. First, using standard higher-order resources alone, we can define, for any property F , a property $F^{\pm N}$ that differs from F only in whether or not it applies to a given entity N . When F does apply to N , $F^{\pm N}$ does not, and vice versa.⁴⁰ We then regiment the metaphysical version

⁴⁰Formally,

$$(\cdot)^{\pm(\cdot)} := \lambda X y z. (Xy \rightarrow (Xz \wedge z \neq y)) \wedge (\neg Xy \rightarrow (Xz \vee z \equiv y)).$$

of quantificational sensitivity as follows:⁴¹

$$\text{Sens}_{\sim} MN^{\sigma} := \exists^{\top} X^{\sigma \rightarrow t} (\text{At}(X)(M) \not\sim \text{At}(X^{\pm N})(M)).$$

Our official definition of metaphysical predicativity is: M is *metaphysically predicative* when it is quantificationally insensitive to itself, i.e., at all domains of the relevant type, M is equivalent to the result of adding or removing M itself from the domain of quantification:

$$\text{Pred}_{\sim} M := \neg \text{Sens}_{\sim} MM$$

Again, $\sim$ here is a placeholder for any equivalence relation. We say that M is *metaphysically impredicative* when it is not metaphysically predicative, i.e., when it *is* quantificationally sensitive to itself.

To get a feel for how this all works, let us explore the notion of metaphysical predicativity that we get when we think of equivalence as coextensiveness. Take the property of being a natural number. In our higher-order language, we can regiment it in a way that matches its definition given in (11), as follows:

$$\text{Nat} := \lambda x. \forall Y (\text{Hereditary} Y \rightarrow Yx),$$

where *Hereditary* is defined as one would expect. Our background logic proves that *Nat* is metaphysically impredicative in the presence of modest assumptions: for example, adequate arithmetical axioms guaranteeing that *Nat* forms a natural number structure (roughly, something isomorphic to a standard model of Peano arithmetic) plus the assumption that something is not a natural number—for concreteness, Julius Caesar. Under these assumptions, we can prove that *Nat* still does not apply to Julius Caesar relative to the domain of properties identical to *Nat*. However, if we remove *Nat* from this domain, the result is the empty domain, relative to which *Nat* vacuously applies to everything, including Julius Caesar.

Contrast this with the set-theoretic intersection operation. Suppose we equip our higher-order language with a stable constant $\in$ for set membership. We require that $\in$ satisfies enough set theory to ensure we can derive every instance of the schema $A \notin A$. We then add another stable constant $\cap$ of type $e \rightarrow e \rightarrow e$ to regiment intersection, requiring that relative to any domain, the existing members of $A \cap B$ be precisely the existing things that belong to both A and B :

$$\blacksquare \forall x (x \in A \cap B \leftrightarrow x \in A \wedge x \in B).$$

Under these assumptions we can prove that $A \cap B$ is metaphysically predicative. Given any first-order domain F , the members of $\text{At}(F)(A \cap B)$ are exactly those F s that belong to both $\text{At}(F)(A)$ and $\text{At}(F)(B)$. Our background set theory proves that $\text{At}(F)(A \cap B)$ is not a member of itself. So, adding or removing $A \cap B$ itself

⁴¹It is important to distinguish between the definition given and the following:

$$\text{Sens}_{\sim}^* MN^{\sigma} := \blacklozenge (M \not\sim \text{At}(E!^{\pm N})(M)).$$

Intuitively, the definition in the main text says that for some domain F , the result of setting F as the domain of quantification within M is not equivalent to the result of setting F plus or minus N as the domain of quantification within M . On the other hand, the definition just given says that, domain-occasionally, M is not equivalent to the result of adding or removing N from the *current* domain of quantification within M . The difference is that, at a domain F , the latter adds or removes not N itself but $\text{At}(F)(N)$, by (Comp) and (App). The two come apart when N is not stable.

from F cannot change what the members of $A \cap B$ are at the resulting domain: $\text{At}(F)(A \cap B)$ and $\text{At}(F^{\pm A \cap B})(A \cap B)$ must have the same members.

My notion of metaphysical predicativity, then, seems to be on the right track: it correctly classifies at least these two paradigmatic cases. I close by discussing its philosophical fruitfulness.

5.3. Metaphysical Predicativism. The sort of predicativist I have sketched in Section 5.2, recall, wants to impose predicativity as a necessary condition on existence. We can straightforwardly do so through the notion of metaphysical predicativity:

Metaphysical Predicativism_~: $\forall x \text{Pred}_{\sim} x$.

This principle is schematic in two ways: different instances are obtained by fixing the type of x and selecting a particular equivalence relation. In what follows, I will call *metaphysical predicativism* the theory that results by adding every instance of the schema above where $\sim$ is coextensiveness and x is of any type.

We have seen that an important application of predicativist theses like Russell's VCP is blocking certain paradoxes. It is fair to wonder whether metaphysical predicativism can do the same. I suggest it can. Metaphysical predicativism can be used as an add-on to solutions to certain paradoxes that work by denying the existence of certain problematic entities. It is used to *explain* why these problematic entities do not exist: they are metaphysically impredicative, and there are no such things as metaphysically impredicative entities. As a proof of concept, I discuss solutions along these lines to two paradoxes: the *Fine–Krämer puzzle of ground* and *Prior's paradox*.

The Fine–Krämer puzzle of ground (Fine, 2010; Krämer, 2013) is a paradox about the relationship between grounding, truth and quantification. Here is a sketch of Krämer's simpler version of the puzzle. Grounding theorists widely hold that, much like true disjunctions are grounded by each of their true disjuncts, true existential generalizations are grounded by each of their true instances (Fine, 2012). If Fx is true and x exists, then the proposition that something is F is grounded by Fx . Formally, with $\prec: t \rightarrow t \rightarrow t$ expressing the grounding relation,

$$P \wedge \text{E!}a \rightarrow P \prec \exists x P[x/a]. \quad (\exists\text{Gr})$$

Grounding theorists also widely hold that grounding is irreflexive:

$$P \not\prec P. \quad (\text{Irr})$$

Consider now the truth-teller proposition, $\exists p p$. Suppose it is true and it exists. Being a true existential generalization, the truth-teller is grounded by each of its true instances as per $(\exists\text{Gr})$. But the truth-teller is its own instance: the truth-teller is $\beta\eta$ -equivalent, hence identical to the proposition that it is true that there is a truth, $(\lambda p.p)(\exists p p)$. So, the truth-teller is self-grounded, contradicting (Irr) .

We can block the argument just given by denying that the truth-teller exists. This move should come with a story about why this is so, lest it seem ad hoc. Both Fine (2010, p. 104) and Krämer (2013, p. 88) flirt with the idea that considerations of predicativity could be brought to bear in telling such a story, but neither develops the idea in detail. Metaphysical predicativism can do just that. The truth-teller is metaphysically impredicative, and there are no such things as metaphysically impredicative propositions.

The truth-teller is metaphysically impredicative under the assumption that it is true. Consider the domain of falsehoods, $\lambda p. \neg p$. The truth-teller is false at this domain, which only contains falsehoods. But adding the truth-teller itself to the domain of falsehoods yields a domain that verifies the truth-teller. For the resulting domain does contain a truth, namely the truth-teller itself. The truth-teller, then, is sensitive to itself, hence metaphysically impredicative. Metaphysical predicativism prohibits its existence, so that $(\exists \text{Gr})$ cannot be used to derive the paradoxical self-grounding claim. The non-existence of the truth-teller is not simply stipulated to block the argument. Under metaphysical predicativism, it is derived from and explained by its metaphysical impredicativity.

Consider next *Prior's paradox* (Prior, 1961). At midnight, Alice sincerely utters the sentence ‘everything I say at midnight is false.’ Intuitively, by doing so Alice succeeds in saying exactly one thing at midnight, namely, that everything she says at midnight is false. Curiously, this cannot be.

Formally, regiment the proposition that everything Alice says at midnight is false as

$$\forall p(Sp \rightarrow \neg p),$$

where S stands for ‘is said by Alice at midnight.’ Call this proposition the *Prior proposition*. Assume the Prior proposition is said by Alice at midnight: $S(\forall p(Sp \rightarrow \neg p))$. Under this assumption, the Prior proposition cannot be true. If it were, i.e., if it were true that everything Alice said at midnight is false, then the Prior proposition itself would be false, being one of the things Alice said at midnight. So the Prior proposition is false, meaning its negation $\exists p(Sp \wedge p)$ is true: there must be some *true* proposition that Alice said at midnight. This true proposition must be different from the Prior proposition, which is false. Alice must have said at least two distinct things at midnight. Nothing in the argument depends on the interpretation of S . As is well known, the paradox can be run with any operator S in place of ‘Alice says at midnight that,’ for instance ‘it seems as though Alice, at midnight, says that’ or ‘Alice, at midnight, fears that.’

The argument just given relies implicitly on the assumption that the Prior proposition exists, which ensures that we can instantiate the universal quantifier in $\forall p(Sp \rightarrow \neg p)$ with the Prior proposition itself. The argument can be blocked by rejecting this assumption.⁴² Solutions that do so have some advantages. In a positive free logic, denying the existence of the Prior proposition allows one to hold onto the intuitive description of what Alice said at midnight: Alice said at midnight that everything she said at midnight is false, and she said only that:

$$S(\forall p(Sp \rightarrow \neg p)) \wedge \forall q(Sq \rightarrow q \equiv \forall p(Sp \rightarrow \neg p)). \quad (14)$$

This is because, even though she did say at midnight that everything she said at midnight is false, she said nothing at all at midnight.

As before, one should tell a story about why $\forall p(Sp \rightarrow \neg p)$ does not exist. It is natural to appeal to impredicativity. The proposition that everything Alice said at midnight is false does intuitively exhibit some form of impredicativity: it seems to

⁴²Bacon et al. (2016) discuss solutions of this sort at length.

quantify over the things that Alice says at midnight, one of which is that proposition itself.⁴³

Metaphysical predicativism vindicates this intuition. In the presence of (14) above, we can prove that the Prior proposition is impredicative. If (14) holds, the Prior proposition must be (vacuously) true and must not exist. But if we add the Prior proposition to the current domain, the result is a domain that falsifies the Prior proposition: the resulting domain does contain a true proposition that Alice said at midnight, namely the Prior proposition itself. The proposed predicativist solution to Prior’s paradox, then, begins by endorsing the intuitive description of what went on at midnight, formalized by (14). It then shows that the Prior proposition cannot exist, and explains its non-existence via its metaphysical impredicativity.

A distinguishing feature of Prior’s paradox is that the Prior proposition is only contingently paradoxical. There is nothing paradoxical in me truly saying that everything that Alice said at midnight is false, for example if all she said at midnight was that snow is green and she did not herself say the Prior proposition at midnight. Metaphysical predicativism can explain this contingency as contingency in the impredicativity of the Prior proposition. There are models of quantificationalism and metaphysical predicativism in which the Prior proposition exists, is true, and not said by Alice at midnight.

I have given but sketches of solutions to the Fine–Krämer puzzle of ground and Prior’s paradox based on metaphysical predicativism. There is a lot more work to be done in fleshing out the details of these solutions and exploring how they compare to alternative approaches. What I hope to have shown with these sketches is that metaphysical predicativity has philosophical potential. Quantificationalism, by providing a framework in which metaphysical predicativity can be formulated, inherits this potential.

6. THE CONSISTENCY OF QUANTIFICATIONALISM

I have postponed discussion of the most important foundational question about quantificationalism: the question of whether it is *consistent*. I did so because it is pointless to inquire into the consistency of a theory that shows no philosophical promise. By now, I hope to have shown that quantificationalism does show philosophical promise, so it is time to turn to the question of consistency.

Here is the short story: unfortunately, the logical framework I sketched for quantificationalism in Section 3 is *not* consistent. However, consistency can be restored by embracing a generalization of that framework where domain specifiers are *genuinely syncategorematic expressions*, in a sense to be explained. This last section is devoted to unpacking the short story.

6.1. Paradox. Recall the distinction between categorematic and syncategorematic expressions. Categorematic expressions are those expressions of a language that can be assigned a type. They include constants, variables, and complex terms that we can form from simpler ones via the term formation rules of the language. Syncategorematic expressions, on the other hand, serve as linguistic markers of

⁴³Russell (1908) blames a very similar paradox to impredicativity. The point that the Prior proposition itself exhibits some kind of impredicativity is made explicitly, e.g., in Kaplan (1995); Bacon et al. (2016); Bacon (2021).

how a given term was constructed. They notate the term formation rules of the language themselves.

I argued towards the end of Section 3.3 that interpretation of domain specifiers in terms of quantificational substitutions, recall, motivates a syncategorematic treatment of domain specifiers. If the correct way of thinking about how domain specifiers compose is given by (Comp), as the substitutional picture suggests, then we should not think that domain specifiers combine with their argument by application. The argument, which I gave informally, can in fact be turned into a *proof* to the effect that, fixing everything else in the framework, domain specifiers cannot be categorematic terms.

Suppose we had treated domain specifiers categorematically. Instead of having a term formation rule for introducing domain specifiers, we would have posited an infinite family of constants $\text{At}_{\sigma\tau}$, each of type $(\sigma \rightarrow t) \rightarrow \tau \rightarrow \tau$. Then a term of the form $\text{At}_{\sigma\tau}(G)(M)$ would be constructed by *applying* the categorematic term $\text{At}_{\sigma\tau}(G)$ to M . So, the requirement that domain specifiers commute with application (App) would predict (omitting subscripts for readability)

$$\text{At}(F)(\text{At}(G)(M)) \equiv \text{At}(F)(\text{At}(G))(\text{At}(F)(M)). \quad (\star)$$

But $(\star)$ is inconsistent in the ambient logic. The reason why is that it imposes severe restrictions on how many stable properties there can be. If we assume that F and G are two stable properties of the same type, then we can use $(\star)$ to show that F and G must be identical. I leave the argument to a footnote.⁴⁴ So, there can be at most one stable property at any given type. At the same time, the ambient logic *proves* the distinctness of at least two stable properties. Truth $(\lambda p.p)$ and negation are the most obvious examples at type $t \rightarrow t$.

We derived $(\star)$ by assuming a categorematic approach to domain specifiers, which the present framework rejects. Unfortunately, the categorematic and syncategorematic approaches to At are less sharply distinct than they might initially appear, meaning the problem just presented arises for my framework as well.

The collapse between the categorematic and syncategorematic approaches to domain specifiers is due to two main features of the ambient logic:

Unrestricted λ -abstraction: given any term M of type τ and any variable x of type σ , we can form the term $\lambda x.M$ of type $\sigma \rightarrow \tau$;

$\beta\eta$ intersubstitutivity: $\beta\eta$ -equivalent terms are intersubstitutable in all contexts.

⁴⁴Assume that F and G are two such properties. Begin with an instance of (Gen):

$$\text{At}(F)(E!) \equiv F.$$

Specifying the G s as the domain in both sides should preserve the identity:

$$\text{At}(G)(\text{At}(F)(E!)) \equiv \text{At}(G)(F).$$

Because F is stable, the right-hand side reduces to F . But, using $(\star)$, we can show that the left-hand side reduces to G . One application of $(\star)$ gives us

$$\text{At}(G)(\text{At}(F)(E!)) \equiv \text{At}(G)(\text{At}(F))(\text{At}(G)(E!))$$

By (Gen) we have $\text{At}(G)(E!) \equiv G$ and by the stability of G we have $G \equiv \text{At}(G)(G)$. So, we may substitute $\text{At}(G)(G)$ for $\text{At}(G)(E!)$ in the equality above, yielding

$$\text{At}(G)(\text{At}(F)(E!)) \equiv \text{At}(G)(\text{At}(F))(\text{At}(G)(G)).$$

Applying $(\star)$ once more to “pull out” the domain specifier $\text{At}(G)$ we see that the right-hand side equals $\text{At}(G)(\text{At}(F)(G))$. Since G is stable, this reduces to G , so we may conclude that $G \equiv F$.

Both are fairly standard assumptions in higher-order logic, though of course neither is without detractors.

Now, consider a term $\text{At}(F)(M)$ constructed via syncategorematic domain specification. Given two variables X and Y whose types match those of F and M , respectively, [Unrestricted \$\lambda\$ -abstraction](#) allows us to form the term $\lambda XY. \text{At}(X)(Y)$. This is a term that can be *applied* to F and M in this order to yield $(\lambda XY. \text{At}(X)(Y))FM$. The latter is $\beta\eta$ -equivalent, hence intersubstitutable with $\text{At}(F)(M)$ by [\$\beta\eta\$ intersubstitutivity](#).

In this sense, the present framework treats domain specifiers as merely *superficially syncategorematic*. In any context, we can always find a categorematic term $\lambda X. \text{At}(F)(X)$ that can be used to *stand in* for the syncategorematic domain specifier $\text{At}(F)$. While we did not set out to treat domain specifiers categorematically, we do have in the language infinitely many categorematic expressions that do exactly the same job that the primitive categorematic constants $\text{At}_{\sigma\tau}$ would.

Consequently, while $(\star)$ itself is not derivable in the present system, a closely related principle is:

$$\text{At}(F)(\text{At}(G)(M)) \equiv \text{At}(F)(\lambda x. \text{At}(G)(x))(\text{At}(F)(M)). \quad (\star\star)$$

The argument used to derive from $(\star)$ that there is at most one stable property at any type can be straightforwardly adapted to show that $(\star\star)$ has the same consequence.

6.2. Genuine Syncategorematicity. Broadly speaking, two kinds of responses are open. We can either revise the underlying framework, or conclude that quantificationalism itself is mistaken.

Because quantificationalism is philosophically promising, it is worth investigating solutions of the former sort rather than hastily concluding that the view itself is mistaken. Abandoning the promising philosophical applications of quantificationalism is a cost, which should be weighed against the costs of revising the underlying framework in such a way as to provide consistent foundations for such applications. I conclude this paper by outlining one such solution and discussing some of its metaphysical implications.

Before I do, I should emphasize that my solution is only one among many possible revisions to the framework one might make to get around the problem just presented, and that I do not intend to argue that my proposal is better than any such alternative. This is a question that cannot be answered before all such alternatives are worked out in reasonable detail, and one has to start somewhere. The present approach is a reasonable first step due to its internal coherence and its natural fit with the suggestive picture of domain specification as a kind of metaphysical substitution.

Here is the key idea. Application and domain specification really do describe distinct ways of combining higher-order entities, which should not be reduced to one another. We linguistically marked this distinction by equipping our language with distinct syncategorematic expressions corresponding to these operations. However, because our framework satisfies [Unrestricted \$\lambda\$ -abstraction](#) and [\$\beta\eta\$ intersubstitutivity](#), this is not a sufficiently strong way of marking the distinction: syncategorematic

domain specifiers have categorematic surrogates that combine with their argument by application. We need a stronger way of marking the distinction.⁴⁵

My preferred approach is to do so by restricting [Unrestricted \$\lambda\$ -abstraction](#) in such a way that categorematic surrogates for domain specifiers, of the form

$$\lambda x.\text{At}(F)(x),$$

cannot be constructed in the first place. The simplest way of achieving this is to prohibit the abstraction of a variable occurring within the scope of a domain specifier from *outside* the scope of that domain specifier. This means that terms like

$$\lambda x.\text{At}(F)(Gx) \qquad \lambda X.\text{At}(F)(Xa)$$

will also be classified as ungrammatical. On the other hand, abstractions of variables falling within the scope of the *first* argument place of *At* from outside the scope of that occurrence of *At*, such as

$$\lambda X.\text{At}(X)(P)$$

remain well-formed.⁴⁶

[\$\beta\eta\$ intersubstitutivity](#) remains true. *Well formed* $\beta\eta$ -equivalent terms remain intersubstitutable in all contexts. The only difference is that, because of the restrictions on λ -abstraction, some terms that used to count as well-formed $\beta\eta$ -equivalents of a given term are no longer well-formed. In particular, $(\lambda x.\text{At}(F)(x))M$ cannot be substituted for $\text{At}(F)(M)$ in any context, but only because ill-formed terms cannot be substituted for well-formed ones, not because $\beta\eta$ -equivalence fails to imply intersubstitutivity.

On this approach, domain specifiers are *genuinely* rather than *superficially* syncategorematic. They are introduced as syncategorematic expressions, and the framework prohibits the construction of categorematic surrogates through λ -abstraction. No other way of uniformly generating such categorematic surrogates exists: in the new system, we can *prove* that many domain specifiers have no categorematic surrogates.⁴⁷

Once we restrict abstraction within the scope of domain specifiers, a couple more changes are in order to obtain an internally coherent framework. First, it seems natural to generalize our restrictions on abstraction to prohibit the abstraction of variables falling within the scope of an occurrence of $\blacksquare$ from outside that occurrence. Terms like

$$\lambda p.\blacksquare p \qquad \lambda x.\blacksquare Fx$$

⁴⁵Other possible “positive” reactions to the problem at hand include (i) somehow restricting ([Comp](#)) or ([App](#)), (ii) rejecting ([Gen](#)), (iii) rejecting [\$\beta\eta\$ intersubstitutivity](#). Note the latter would also require denying the much weaker principle that $\beta\eta$ -equivalent terms are *coextensive*.

⁴⁶Thus domain specification does not have a categorematic counterpart, though for each particular M there does exist a categorematic term whose semantic function is to take a property and return the result of specifying that property as the domain of quantification within M .

⁴⁷Another approach I consider continuous with the present one is to introduce a special base type i of metaphysical substitutions and regard domain specifiers as terms of type i . Domain specifiers would thus be categorematic but not “applicative” expressions: but we would need a primitive syncategorematic operation allowing us to combine a term of type i with one of some type σ other than i in order to form terms such as $\text{At}(F)(M)$. This operation would have to be regarded as genuinely syncategorematic, on pain of reintroducing surrogates of domain specifiers of functional types $\sigma \rightarrow \sigma$, which is what led to our problem in the first place. I am intrigued by this approach, but I have not investigated whether it can be consistently worked out.

are thus ill-formed. Intuitively, the term $\blacksquare p$ is equivalent to an infinite conjunction saying that p is true at every (in the classical sense) domain of type $\sigma \rightarrow t$, and at every domain of type $\tau \rightarrow t$, and so on for all types.⁴⁸ Were infinite conjunctions of this form expressible in our language, abstracting the variable p from outside the scope of the infinite conjunction would be prohibited, hence we should regard the abstraction $\lambda p.\blacksquare p$ as ill-formed.

Second, following standard conventions in higher-order logic I have been treating quantifiers as categorematic constants. The notation $\forall x^\sigma P$ abbreviates the term $\forall_\sigma \lambda x^\sigma.P$. This means that we cannot express the claim that P holds for every x unless the abstraction $\lambda x.P$ is well formed. Now, one of the generalizations we certainly want to express is quantificationalism itself:

$$\exists p(\blacklozenge p \wedge \blacklozenge \neg p).$$

but the abstraction $\lambda p.\blacklozenge p \wedge \blacklozenge \neg p$ is ill-formed.

The solution is to treat quantifiers as variable-binding syncategorematic expressions, the way we are used to doing in first-order logic. Instead of having infinitely many constants $\forall_\sigma$ each of type $(\sigma \rightarrow t) \rightarrow t$, we add an extra term formation rule to our language: if P is a term of type t and x a variable of any type, then $\forall x P$ is a term of type t . Thus, we allow syncategorematic quantifiers to bind variables in positions where λ cannot.

It is important to note that the reason for treating quantifiers syncategorematically is very different from the reason for treating domain specifiers and $\blacksquare$ syncategorematically. The latter choice is meant to avoid the collapse of two operations we want to sharply distinguish: application and domain specification. The former choice is motivated by a need for expressive power.

Thus, the philosophical reasons given to restrict abstraction within the scope of domain specifiers and $\blacksquare$ do not motivate an analogous restriction for quantifiers. In particular, there is no reason to prohibit the formation of categorematic surrogates for syncategorematic quantifiers. These can be defined as follows:

$$\hat{\forall}_\sigma := \lambda X^{\sigma \rightarrow t}.\forall y Xy \qquad \hat{\exists}_\sigma := \lambda X^{\sigma \rightarrow t}.\exists y Xy.$$

These expressions stand for higher-order entities that are in perfectly good standing. They simply will not be as powerful as syncategorematic quantifiers, as they can only generalize in positions where λ can bind.

Being categorematic expressions, $\hat{\forall}$ and $\hat{\exists}$ can occur as arguments of domain specifiers. The principles about quantifiers of the form $\forall^F$ discussed in Section 3.3, should now be read as principles about $\text{At}(F)(\hat{\forall})$: with syncategorematic quantifiers, $\text{At}(F)(\forall)$ is not even well-formed. The ambient logic still constrains the behavior of syncategorematic quantifiers, though not exactly in the same way as it does for categorematic ones.⁴⁹ Besides occasionally replacing $\forall$ with $\hat{\forall}$, no logical principle formulated in the earlier sections of the paper needs to be revised. While I reserved discussion of limitations on λ -abstraction to the end of the paper, I took care not to violate them throughout.

⁴⁸More precisely:

$$\forall_{\sigma \rightarrow t}^\top X \text{At}(X)(p) \wedge \forall_{\tau \rightarrow t}^\top Y \text{At}(Y)(p) \wedge \forall_{\rho \rightarrow t}^\top Z \text{At}(Z)(p) \wedge \dots$$

and so on for all types.

⁴⁹In particular, when x is not abstractable in P , $\text{At}(F)(\exists x P)$ can be inferred from P and $\text{At}(F)(\text{E! } x)$, though not in general from P and Fx .

The framework that results from making these adjustments—restrictions on λ -abstraction and syncategorematic treatment of quantifiers—supports a consistent logic of domain specifiers. The paradox from Section 6.1 cannot be reproduced. Consistency can be proved by establishing that the system is sound with respect to a class of structures I call *quantificational substitution structures*, which formally regiment the suggestive interpretation of domain specifiers in terms of metaphysical substitutions.⁵⁰

6.3. Consequences of Genuine Syncategorematicity. Embracing genuine syncategorematicity has significant implications for the metaphysics of identity. The resulting picture is unusual, but ultimately illuminating. This last section is devoted to fleshing out this picture.

I have provisionally been treating the typed relations $\equiv_\sigma$ as higher-order identity relations, cautioning that this claim would have to eventually be revised. The core inferential role of the relations $\equiv_\sigma$ is captured by a common schematic regimentation of Leibniz’s Law in higher-order free logics, which we may call *applicative Leibniz’s Law*:

$$M \equiv N \rightarrow FM \rightarrow FN. \quad (\text{ALL})$$

An instance of (ALL) says that given $M \equiv N$, if F *applies* to M , then F also *applies* to N . There is another way of formulating Leibniz’s Law, which might be called *substitutional Leibniz’s Law*:

$$M \equiv N \rightarrow P[M/x] \rightarrow P[N/x]. \quad (\text{SLL})$$

An instance of (SLL) says that if $M \equiv N$, then M is *substitutable* for N in P *salva veritate*.

In standard higher-order frameworks that satisfy both *Unrestricted λ -abstraction* and *$\beta\eta$ intersubstitutivity*, (ALL) and (SLL) are equivalent. Clearly, any instance of (ALL) is an instance of (SLL). Conversely, a claim of the form $P[M/x]$ can always be rewritten as the result of *applying* the property $\lambda x.P$ to M , so that any instance of (SLL) is $\beta\eta$ -equivalent to an instance of (ALL).

However, in a framework that rejects *Unrestricted λ -abstraction* such as ours, there is no guarantee that $\lambda x.P$ is always going to be well formed. For example, from $P = Q$ and $\text{At}(F)(P)$, we cannot use (ALL) to conclude that $\text{At}(F)(Q)$, because $\lambda p.\text{At}(F)(p)$ is ill-formed. What is more, in our newly revised framework, there is no purely logical argument from $P = Q$ and $\text{At}(F)(P)$ to $\text{At}(F)(Q)$. Counterexamples to (SLL) are perfectly consistent, and indeed quite widespread in the sort of models I use to establish consistency.

One could react to this result by concluding that *identity* does not guarantee intersubstitutivity in all contexts *salva veritate*.⁵¹ That is not my reaction. I take it as granted that whatever identity is, identical things ought to be intersubstitutable in all contexts *salva veritate*, and on that basis conclude that the constants $\equiv_\sigma$ do not express higher-order identity relations. Officially, I call the relations expressed by the constants $\equiv_\sigma$ *applicative indiscernibility* relations. ‘Applicative indiscernibility’ is a theoretical term whose meaning is given by its theoretical role, consisting of (ALL) and a reflexivity axiom.

⁵⁰I give some detail in Section A, and a lot more detail in the companion paper [REDACTED].

⁵¹This is the reaction of those who countenance *opacity*: failures of substitutivity of identicals in some contexts, such as attitude reports. See especially Bacon and Russell (2019).

Identity and applicative indiscernibility are relations that appear to coincide from the point of view of standard higher-order framework with [Unrestricted \$\lambda\$ -abstraction](#), but are revealed to be distinct once the latter assumption is weakened. The theoretical role defining applicative indiscernibility happens to pick out identity in contexts where [Unrestricted \$\lambda\$ -abstraction](#) is assumed, but need not do so if the assumption is dropped. The dialectical situation is analogous to that which arises between the classical and the free logicians with respect to the relations of identity and *Leibniz equivalence* respectively. In classical higher-order logic, it is common to *define* identity as the relation two entities bear when they share all their higher-order properties:

$$\lambda xy. \forall Z (Zx \leftrightarrow Zy).$$

Leibniz equivalence can easily be shown to satisfy a substitutional formulation of Leibniz’s law in classical higher-order logics with [Unrestricted \$\lambda\$ -abstraction](#). It suffices to point out that, given classical logic, Leibniz equivalence satisfies an applicative version of Leibniz’s law, and we know that given [Unrestricted \$\lambda\$ -abstraction](#) the difference between applicative and substitutional versions of Leibniz’s law is merely notational.

In a free logic, however, universal generalizations fail to imply all their instances. So, two Leibniz equivalent entities may differ in whether they are F , for F a particular predicate. They must therefore be distinct. Leibniz equivalence and identity are two relations that are taken to coincide by the classical logician, but that the free logician must distinguish.⁵²

A counter-example to [\(SLL\)](#) involves two entities that can only be distinguished by means of genuinely syncategorematic expressions. There are principled quantificationalist metaphysical pictures on which violations of [\(SLL\)](#) are expected and quite pervasive.

An example is a picture in which metaphysical necessity does not imply domain-invariable truth: certain domains of quantification are not metaphysically possible. There are several ways of motivating this thought, but an intriguing one comes from the literature on higher-order contingentism ([Fritz and Goodman, 2016](#)). Higher-order contingentists typically hold that as a matter of metaphysical necessity, higher-order entities can only exist if all the individuals they are about exist as well. The proposition that Gavin Newsom is tall, for example, is about Gavin Newsom and depends on him for its existence in this sense. On the other hand, as we have seen, the present framework holds that existence properties at different types are freely recombinable: what exists at one type places no restrictions whatsoever on what exists at other types. So, on this picture, metaphysical necessity is *narrower* than domain-invariable truth: certain claims linking what exists at higher types with what exists at the type of individuals will be metaphysically necessary but not domain-invariably true.

⁵²Now, if a free logician is in the position to *define* a classical quantifier, then they can reduce identity to *classical* Leibniz equivalence, which is the relation of sharing all higher-order properties in the classical sense of “all.” The present framework uses a primitive free quantifier and has the resources to define classical quantifiers, but we cannot reduce identity to classical Leibniz equivalence either because the latter does not in general satisfy a substitutional version of Leibniz’s law in the framework. The system does, however, prove that classical Leibniz equivalence and applicative indiscernibility are logically equivalent: classical Leibniz equivalence is defined using *stable* quantifiers, which makes the relation itself stable.

An example is *the proposition that Gavin Newsom is tall exists only if Gavin Newsom exists*—call this proposition P —is metaphysically necessary. Since the Boolean tautology $\top$ is also metaphysically necessary, if metaphysically necessary equivalence implies applicative indiscernibility, then P and $\top$ are applicatively indiscernible. But they cannot be identical. There are domains at which, say, the proposition that Gavin Newsom is tall exists but Gavin Newsom does not. Any such domain verifies $\top$ but must falsify P .

Another very general thesis that generates pervasive counter-examples to (SLL) is a comprehension principle for stable entities. Roughly, such a principle asserts that any entity is applicatively indiscernible from a stable entity called its *stabilizer*. Consider, for example, the non-stable property of *loving someone*. Its stabilizer is a stable property whose extension at any domain is those who love someone *in the current domain*, rigidly specified, which cannot be distinguished from the property of loving someone except by genuinely syncategorematic expressions. The function of stabilizers is somewhat similar to that of actuality operators in modal logic (Crossley and Humberstone, 1977). The property of *actually loving someone* is not the same as the property of loving someone, but the two can only be distinguished via modal vocabulary. An entity and its stabilizer always generate a counter-example to (SLL), since the former but not the latter can be “moved” by a domain specifier.

If the quantificationalist embraces violations of (SLL), how should they theorize about identity? Perhaps we could introduce additional typed constants into the language standing for higher-order identity relations $=_\sigma$ and require them to satisfy appropriate substitutional versions of Leibniz’s Law. Unfortunately, this is not possible, since doing so would reintroduce categorematic surrogates of domain specifiers into the system: one could simply write

$$\lambda x. \text{At}(F)(y) \wedge x = y$$

instead of $\lambda x. \text{At}(F)(x)$. In the above, x does not occur free within the scope of a domain specifier, so the term is well formed. But essentially the same argument from Section 6.1 can run using the new term.

This example suggests that if we want to add identity to the language, it should also be a genuinely syncategorematic expression. Happily, the quantificationalist has a story to tell about why we should *expect* identity to be genuinely syncategorematic. In a slogan: *identity is domain-invariable applicative indiscernibility*. Since identity is *defined* in terms of domain-invariable truth this way, its genuine syncategorematicity is no surprise. It is simply another consequence of the genuine syncategorematicity of domain specifiers.

What I mean by this slogan is that the quantificationalist can consistently endorse a substitutional version of Leibniz’s law for the composition of applicative indiscernibility with the operator $\blacksquare$:

$$\blacksquare(M \equiv N) \rightarrow P[M/x] \rightarrow P[N/x], \quad (\text{SLL}_{\blacksquare})$$

where x must be free for both M and N in P .⁵³ Entities that are domain-invariably applicatively indiscernible are intersubstitutable in all contexts.

We know that applicatively indiscernible entities need not be domain-invariably applicatively indiscernible. This fact is exactly what allowed us to generate counter-examples to (SLL). On the higher-order contingentist picture sketched earlier, for example, $\top$ and *the proposition that Gavin Newsom is tall exists only if Gavin Newsom exists* were in fact applicatively indiscernible, but became applicatively discernible once we specified the domain to be one where Gavin Newsom does not exist but the proposition that he is tall does. Relative to such a domain, the two propositions are discernible via the property of being true. What (SLL_■) tells us is that domain specification is the *only* way we can distinguish between two applicatively indiscernible entities. If two entities are not just in fact applicatively indiscernible, but remain so no matter what we specify the domain to be, then there is no way to distinguish between them whatsoever.

Distinguishing between applicative indiscernibility and identity is an undoubtedly revisionary move. But the distinction is principled: it is downstream of the philosophically motivated commitment to the genuine syncategorematicity of domain specifiers. Moreover, it is illuminating in that it makes room for a distinctly quantificationalist understanding of identity itself.

CONCLUSION

I take myself to have shown that quantificationalism is both philosophically promising and considerably more insidious than it might initially appear. The seemingly simple idea that truth can vary across domains in much the same way it can vary across time and worlds has a surprisingly wide range of potential applications, of which we have barely begun to scratch the surface. At the same time, the project of making precise sense of domain-occasional truth faces challenges that simply do not arise in the case of temporary and contingent truth. I proposed one way to meet these challenges, which while principled is certainly not without costs.

Quantificationalism should be judged by weighing the benefits of its philosophical applications against the costs incurred in consistently formulating it. I do not think this is a judgment we are in a position to make at present. But I do believe that the prospects are promising enough to warrant further investigation.

REFERENCES

Bacon, A. [2018]. Tense and Relativity. *Noûs*, 52(3):667–696.

⁵³Officially, I treat (SLL_■) as an axiom schema, but it is worth noting that it can be derived as a theorem of the system obtained by replacing (SLL_■) with appropriate principles of propositional granularity. An example is an axiom schema to the effect that propositions are applicatively indiscernible when domain-invariably materially equivalent:

$$\blacksquare(p \leftrightarrow q) \rightarrow p \equiv q. \quad (\text{QFreg})$$

Alternatively, we could replace (SLL_■) with an inference rule, derivable from (QFreg), asserting that if we *prove* that p and q are materially equivalent, then they must be applicatively indiscernible:

$$\text{from } \vdash p \leftrightarrow q \text{ infer } p \equiv q. \quad (\text{Equiv})$$

The reason is that (Equiv), and therefore (QFreg), ensures that $P \equiv Q$ implies both $\forall x P \equiv \forall x Q$ and $\blacksquare P \equiv \blacksquare Q$. These claims, notice, cannot be straightforwardly derived from (ALL) because both the quantifier and $\blacksquare$ are syncategorematic terms. Once we have them, we can show that $\blacksquare(A \equiv B)$ implies $\blacksquare(M[A/x] \equiv M[B/x])$ by induction on the structure of M .

- [2019]. Substitution Structures. *Journal of Philosophical Logic*, 48(6):1017–1075.
- [2020]. Logical Combinatorialism. *Philosophical Review*, 129(4):537–589.
- [2021]. Opacity and Paradox. In C. Nicolai and J. Stern, eds., *Modes of Truth: The Unified Approach to Truth, Modality, and Paradox*, pp. 231–265. New York: Routledge.
- [2023a]. *A Philosophical Introduction to Higher-order Logics*. Oxford University Press.
- [2023b]. A Theory of Structured Propositions. *Philosophical Review*, 132(2):173–238.
- [2024]. Mathematical Modality: An Investigation in Higher-order Logic. *Journal of Philosophical Logic*, 53(1):131–179.
- Bacon, A., Hawthorne, J., and Uzquiano, G. [2016]. Higher-Order Free Logic and the Prior-Kaplan Paradox. *Canadian Journal of Philosophy*, 46(4-5):493–541.
- Bacon, A. and Russell, J. S. [2019]. The Logic of Opacity. *Philosophy and Phenomenological Research*, 99(1):81–114.
- Brogaard, B. [2006]. Tensed Relations. *Analysis*, 66(3):194–202.
- Buridan, J. [1976]. *Iohannis Buridani Tractatus de Consequentibus*, *Philosophes médiévaux*, vol. 16. Louvain: Publications universitaires. Ed. Hubert Hubien.
- Chihara, C. [1973]. *Ontology and the Vicious-Circle Principle*. Cornell University Press.
- Crossley, J. N. and Humberstone, L. [1977]. The logic of “actually”. *Reports on Mathematical Logic*, 8:11–29.
- Divers, J. [1999]. A Genuine Realist Theory of Advanced Modalizing. *Mind*, 108(430):217–239.
- [2014]. The Modal Status of the Lewisian Analysis of Modality. *Mind*, 123(491):861–872.
- Divers, J. and Parry, J. J. [2018]. Advanced modalizing de dicto and de re. *Analysis*, 78(3):415–425.
- Dorr, C. [2014]. Quantifier Variance and the Collapse Theorems. *The Monist*, 97(4):503–570.
- [2016]. To Be F Is To Be G. *Philosophical Perspectives*, 30(1):39–134.
- Dorr, C. and Goodman, J. [2020]. Diamonds Are Forever. *Noûs*, 54(3):632–665.
- Feferman, S. [2005]. Predicativity. In S. Shapiro, ed., *Oxford Handbook of Philosophy of Mathematics and Logic*, pp. 590–624. Oxford: Oxford University Press.
- Feferman, S. and Hellman, G. [2021]. Challenges to Predicative Foundations of Arithmetic. In *Mathematics and Its Logics: Philosophical Essays*, pp. 117–138. Cambridge: Cambridge University Press.
- Fine, K. [1977]. Properties, Propositions and Sets. *Journal of Philosophical Logic*, 6(1):135–191.
- [1994]. Essence and Modality. *Philosophical Perspectives*, 8:1–16.
- [2005]. Tense and reality. In *Modality and Tense: Philosophical Papers*, pp. 261–320. Oxford: Oxford University Press.
- [2006a]. The reality of tense. *Synthese*, 150(3):399–414.
- [2006b]. Relatively unrestricted quantification. In A. Rayo and G. Uzquiano, eds., *Absolute Generality*, pp. 20–44. Oxford University Press.
- [2010]. Some Puzzles of Ground. *Notre Dame Journal of Formal Logic*, 51(1):97–118.

- [2012]. Guide to ground. In F. Correia and B. Schnieder, eds., *Metaphysical Grounding*, pp. 37–80. Cambridge University Press.
- [2016]. Williamson on Fine on Prior on the reduction of possibilist discourse. *Canadian Journal of Philosophy*, 46(4-5):548–570.
- Fritz, P. and Goodman, J. [2016]. Higher-Order Contingentism, Part 1: Closure and Generation. *Journal of Philosophical Logic*, 45(6):645–695.
- Fritz, P., Lederman, H., and Uzquiano, G. [2021]. Closed Structure. *Journal of Philosophical Logic*, 50(6):1249–1291.
- Gale, R. M. [1968]. *The Language of Time*. New York: Humanities Press.
- Goodman, J. [2017]. Reality is Not Structured. *Analysis*, 77(1):43–53.
- Goodman, N. [1951]. *The Structure of Appearance*. Cambridge, MA, USA: Harvard University Press.
- Gödel, K. [1944]. Russell’s Mathematical Logic. In P. A. Schilpp, ed., *The Philosophy of Bertrand Russell*, pp. 123–153. Evanston and Chicago: Northwestern University.
- Harris, J. H. [1982]. What’s so Logical About the “Logical” Axioms? *Studia Logica*, 41(2-3):159–171.
- Hawthorne, J. and Cortens, A. [1995]. Towards Ontological Nihilism. *Philosophical Studies*, 79(2):143–165.
- Hirsch, E. [2002]. Quantifier variance and realism. *Philosophical Issues*, 12:51–73.
- [2010]. *Quantifier Variance and Realism: Essays in Metaontology*. New York: Oxford University Press.
- Hoeksema, J. [1987]. The logic of exception. In *Proceedings of the Fourth Eastern States Conference on Linguistics (ESCOL 4)*, pp. 100–113.
- Kaplan, D. [1989]. Demonstratives: An Essay on the Semantics, Logic, Metaphysics and Epistemology of Demonstratives and Other Indexicals. In J. Almog, J. Perry, and H. Wettstein, eds., *Themes From Kaplan*, pp. 481–563. Oxford University Press.
- [1995]. A Problem in Possible Worlds Semantics. In W. Sinnott-Armstrong, D. Raffman, and N. Asher, eds., *Modality, Morality and Belief: Essays in Honor of Ruth Barcan Marcus*, pp. 41–52. Cambridge University Press.
- Kretzmann, N. [1982]. Syncategoremata, exponibilia, sophismata. In N. Kretzmann, A. Kenny, and J. Pinborg, eds., *The Cambridge History of Later Medieval Philosophy*, pp. 211–245. Cambridge: Cambridge University Press.
- Krämer, S. [2013]. A Simpler Puzzle of Ground. *Thought: A Journal of Philosophy*, 2(2):85–89.
- Lewis, D. K. [1968]. Counterpart Theory and Quantified Modal Logic. *Journal of Philosophy*, 65(5):113–126.
- [1980]. Index, Context, and Content. In S. Kanger and S. Öhman, eds., *Philosophy and Grammar*, pp. 79–100. Reidel.
- [1986]. *On the Plurality of Worlds*. Oxford: Blackwell.
- Linnebo, Ø. [2013]. The potential hierarchy of sets. *The Review of Symbolic Logic*, 6(2):205–228.
- [2018]. *Thin Objects*. Oxford University Press.
- McTaggart, J. E. [1908]. The Unreality of Time. *Mind*, 17(68):457–474.
- Mellor, D. H. [1981]. *Real Time*. Cambridge University Press.
- Myhill, J. [1958]. Problems Arising in the Formalization of Intensional Logic. *Logique et Analyse*, 1(1):78–83.

- Prior, A. N. [1957]. *Time and Modality*. Westport, Conn.: Oxford University Press.
- [1961]. On a Family of Paradoxes. *Notre Dame Journal of Formal Logic*, 2(1):16–32.
- [1967]. *Past, Present and Future*. Oxford, GB: Oxford University Press.
- Prior, A. N. and Fine, K. [1977]. *Worlds, Times and Selves*. Duckworth.
- Quine, W. V. [1953]. Three Grades of Modal Involvement. *Proceedings of the XIth International Congress of Philosophy*, 14:65–81.
- Quine, W. V. O. [1960]. *Word and Object*. Cambridge, MA, USA: MIT Press.
- Richard, M. [1981]. Temporalism and Eternalism. *Philosophical Studies*, 39(1):1–13.
- Russell, B. [1903]. *Principles of Mathematics*. Cambridge: Cambridge University Press.
- [1908]. Mathematical Logic as Based on the Theory of Types. *American Journal of Mathematics*, 30(3):222–262.
- [1915]. On the Experience of Time. *The Monist*, 25(2):212–233.
- Russell, B., ed. [1940]. *The Philosophy of Logical Atomism*. Routledge.
- Russell, B. [1973]. *Essays in Analysis*. London, England: Allen & Unwin.
- Schaffer, J. [2012]. Necessitarian Propositions. *Synthese*, 189(1):119–162.
- Sider, T. [2001]. *Four Dimensionalism: An Ontology of Persistence and Time*. Oxford University Press.
- [2009]. Ontological Realism. In D. J. Chalmers, D. Manley, and R. Wasserman, eds., *Metametaphysics: New Essays on the Foundations of Ontology*, pp. 384–423. Oxford University Press.
- [2011]. *Writing the Book of the World*. New York: Oxford University Press.
- Skolem, T. [1967]. Some Remarks on Axiomatized Set Theory. In J. van Heijenoort, ed., *From Frege to Gödel: a Sourcebook in Mathematical Logic*, pp. 290–301. Harvard University Press.
- Spinoza, B. d. [1985]. *The Collected Works of Spinoza*. Princeton, N.J.: Princeton University Press. Ed. E. M. Curley.
- Stanley, J. [2002]. Nominal restriction. In G. Preyer and G. Peter, eds., *Logical Form and Language*, pp. 365–388. Oxford University Press.
- Stanley, J. and Szabó, Z. G. [2000]. On Quantifier Domain Restriction. *Mind and Language*, 15(2-3):219–261.
- Stojanovic, I. [2012]. Domain-sensitivity. *Synthese*, 184(2):137–155.
- Uzquiano, G. [2015]. A Neglected Resolution of Russell’s Paradox of Propositions. *The Review of Symbolic Logic*, 8(2):328–344.
- [2023]. Ramified structure. *Philosophical Studies*, 180(5-6):1651–1674.
- van Inwagen, P. [1990]. *Material Beings*. Cornell University Press.
- Von Stechow, K. [1993]. Exceptive constructions. *Natural Language Semantics*, 1(2):123–148.
- Westerstahl, D. [1985]. Determiners and context sets. In J. van Benthem and A. ter Meulen, eds., *Generalized Quantifiers in Natural Language*, pp. 45–71. Dordrecht: Foris.
- Weyl, H. [1994]. *The Continuum: A Critical Examination of the Foundation of Analysis*. New York: Dover Publications.
- Whitehead, A. N. and Russell, B. [1910]. *Principia Mathematica*. Cambridge University Press.
- Williamson, T. [1988]. Equivocation and Existence. *Proceedings of the Aristotelian Society*, 88(1):109–128.

- [2003]. Everything. *Philosophical Perspectives*, 17(1):415–465.
- [2013a]. How Deep is the Distinction Between A Priori and A Posteriori Knowledge? In A. Casullo and J. C. Thurrow, eds., *The A Priori in Philosophy*, pp. 291–312. Oxford University Press.
- [2013b]. *Modal Logic as Metaphysics*. Oxford University Press.
- Wittgenstein, L. [1921]. *Tractatus Logico-Philosophicus* (Trans. Pears and McGuinness). New York: Routledge. Ed. Luciano Bazzocchi and P. M. S. Hacker.
- Yagisawa, T. [2010]. *Worlds and Individuals, Possible and Otherwise*. Oxford: Oxford University Press.
- Yli-Vakkuri, J., Hawthorne, J., and Fritz, P. [2019]. Operator Arguments Revisited. *Philosophical Studies*, 176(11):2933–2959.
- Zimmerman, D. W. [2005]. The A-Theory of Time, the B-Theory of Time, and ‘Taking Tense Seriously’. *Dialectica*, 59(4):401–457.

APPENDIX A. TECHNICAL DETAILS

This appendix states the core of the formal framework as compactly as possible. Proofs and further discussion can be found in [REDACTED].

A.1. The Language $\mathcal{L}_+$. Types are generated from e and t by $\rightarrow$. A *nice type sequence* is a finite, non-repeating sequence of types $\bar{\sigma}$; we write $\sigma \in \bar{\sigma}$ when σ occurs in $\bar{\sigma}$. Fix a typed supply of variables Var^σ and a signature Σ extending the logical signature Λ , which contains $\wedge : t \rightarrow t \rightarrow t$, $\neg : t \rightarrow t$, and, for each type σ , an existence predicate $E!_\sigma : \sigma \rightarrow t$ and an applicative indiscernibility constant $\equiv_\sigma : \sigma \rightarrow \sigma \rightarrow t$.

The terms of the full language $\mathcal{L}(\Sigma)$ and their types are given by the following clauses, where $M : \sigma$ means that M is a term of type σ :

- (15) $x \cdot \sigma$ for $x \in Var^\sigma$, and $C \cdot \sigma$ for $C \in \Sigma^\sigma$;
- (16) $MN \cdot \tau$ whenever $M \cdot \sigma \rightarrow \tau$ and $N \cdot \sigma$;
- (17) $\lambda x.M \cdot \sigma \rightarrow \tau$ whenever $M \cdot \tau$ and $x \in Var^\sigma$;
- (18) $\forall xP \cdot t$ and $\exists xP \cdot t$ whenever $P \cdot t$ and x is a variable of any type;
- (19) $At(\bar{F})(M) \cdot \tau$ whenever $M \cdot \tau$ and $\bar{F} = (F_\sigma)_{\sigma \in \bar{\sigma}}$ is a *nice term sequence*, i.e., $\bar{\sigma}$ is a nice type sequence and $F_\sigma \cdot \sigma \rightarrow t$ for each $\sigma \in \bar{\sigma}$;
- (20) $\blacksquare P \cdot t$ whenever $P \cdot t$.

The domain specifiers $At(F)$ of the main text abbreviate $At(\bar{F})$ for the unit sequence $\bar{F} = (F)$. Quantifiers, domain specifiers, and $\blacksquare$ are all syncategorematic. We write $FV_\sigma(M)$ for the free variables of type σ in M , and $\blacklozenge P$ for $\neg \blacksquare \neg P$. The categorematic quantifiers are defined by $\hat{\forall}_\sigma := \lambda X^{\sigma \rightarrow t}.\forall y.Xy$ and $\hat{\exists}_\sigma := \lambda X^{\sigma \rightarrow t}.\exists y.Xy$; the other Boolean connectives, $\top$ and $\perp$, and their pointwise liftings to higher types are defined as usual.

The set $AV_\sigma(M)$ of variables of type σ that are *abstractable* in M is defined by:

$$\begin{aligned}
 AV_\sigma(x) &= AV_\sigma(C) := Var^\sigma \\
 AV_\sigma(MN) &:= AV_\sigma(M) \cap AV_\sigma(N) \\
 AV_\sigma(\lambda x.M) &:= AV_\sigma(M) \\
 AV_\sigma(\forall xP) &= AV_\sigma(\exists xP) := AV_\sigma(P) \\
 AV_\sigma(At(\bar{F})(M)) &:= \bigcap_{\tau \in \bar{\sigma}} AV_\sigma(F_\tau) \setminus FV_\sigma(M) \\
 AV_\sigma(\blacksquare P) &:= Var^\sigma \setminus FV_\sigma(P)
 \end{aligned}$$

Thus a variable is abstractable in M unless it occurs free within the second argument place of some occurrence of At , or within the scope of some occurrence of $\blacksquare$. The language $\mathcal{L}_\dagger(\Sigma)$ is defined like $\mathcal{L}(\Sigma)$, except that clause 17 is replaced by:

3*. $\lambda x.M : \sigma \rightarrow \tau$ whenever $M : \tau$ and $x \in \text{AV}_\sigma(M)$.

For example,

$$\lambda x.\text{At}(\bar{F})(x), \quad \lambda X.\text{At}(\bar{F})(Xa), \quad \lambda p.\blacksquare p$$

are not $\mathcal{L}_\dagger$ -terms, while $\lambda X.\text{At}(X)(P)$ is, provided X is not free in P . Quantifiers are not subject to any restriction: $\exists p(p \wedge \blacklozenge \neg p)$ is a $\mathcal{L}_\dagger$ -term. We write $\mathcal{L}_\dagger$ for $\mathcal{L}_\dagger(\Lambda)$.

A.2. The Logic $\mathbf{Q}_\dagger$. Given nice term sequences $\bar{F}$ and $\bar{G}$, indexed by $\bar{\sigma}$ and $\bar{\tau}$ respectively, their *merge* $\bar{F} \bullet \bar{G}$ is the nice term sequence indexed by the types occurring in either $\bar{\sigma}$ or $\bar{\tau}$, with

$$(\bar{F} \bullet \bar{G})_\rho := \begin{cases} \text{At}(\bar{F})(G_\rho) & \text{if } \rho \in \bar{\tau}, \\ F_\rho & \text{if } \rho \in \bar{\sigma} \text{ but } \rho \notin \bar{\tau}. \end{cases}$$

A *combinator* is a closed term built from variables using only application and λ -abstraction. The logic $\mathbf{Q}_\dagger$ is the set of $\mathcal{L}_\dagger$ -terms $P : t$ such that $\vdash P$, where $\vdash$ is the least relation between sets of $\mathcal{L}_\dagger$ -formulas and $\mathcal{L}_\dagger$ -formulas satisfying the conditions in Figures 2 to 4. All schemas are restricted to $\mathcal{L}_\dagger$ -terms of the appropriate types, and a, M, N are assumed to be free for the relevant variables.

$\vdash P$	P a classical tautology	(Taut)
$\vdash \forall x \text{E}!x$		($\forall\text{E}!$)
$\vdash \text{E}!a \rightarrow (\forall x P \rightarrow P[a/x])$		(FrUI)
$\vdash \forall x(P \rightarrow Q) \rightarrow (\forall x P \rightarrow \forall x Q)$		(Norm)
$\vdash \forall x P \leftrightarrow \neg \exists x \neg P$		(Dual)
$\vdash M \equiv M$		(Refl)
$\vdash M \equiv N \rightarrow (FM \rightarrow FN)$		(ALL)
$\vdash M \equiv N$	M, N $\beta\eta$ -equivalent	($\beta\eta$)
$P, P \rightarrow Q \vdash Q$		(MP)
If $\Gamma \vdash Q$ then $\Gamma \vdash \forall x Q$	x not free in Γ	(UG)

FIGURE 2. Free higher-order logic.

In (Surj), $\bar{\sigma}$ ranges over nice type sequences and $\bar{X} : \bar{\sigma}$ means that $\bar{X}$ is a sequence of distinct variables indexed by $\bar{\sigma}$. The rule is infinitary, since it generalizes over all nice type sequences. Where $\bar{F}$ is a sequence of existence predicates, it reduces by (Id) to the simpler rule mentioned in Section 3: if $\text{At}(\bar{X})(P)$ is derivable for arbitrary $\bar{X}$, then so is $\blacksquare P$. In $\mathbf{Q}_\dagger$, $\blacksquare$ is a normal modality satisfying T, 4, and closed under necessitation. Identity is defined by $M = N := \blacksquare(M \equiv N)$, which satisfies substitutional Leibniz's Law by (SLL $\blacksquare$).

$\vdash \text{At}(\bar{E}!)(M) \equiv M$	$\bar{E}! = (E!_{\sigma})_{\sigma \in \bar{\sigma}}$	(Id)
$\vdash \text{At}(\bar{F})(MN) \equiv \text{At}(\bar{F})(M)(\text{At}(\bar{F})(N))$		(App)
$\vdash \text{At}(\bar{F})(\text{At}(\bar{G})(M)) \equiv \text{At}(\bar{F} \bullet \bar{G})(M)$		(Comp)
$\vdash \text{At}(\bar{F})(M) \equiv \text{At}(\bar{G})(M)$	$\bar{G}$ a permutation of $\bar{F}$	(Perm)
$\vdash \text{At}(\bar{F})(E!_{\sigma}) \equiv F_{\sigma}$	$\sigma \in \bar{\sigma}$	(Gen)
$\vdash \text{At}(\bar{F})(E!_{\sigma}) \equiv E!_{\sigma}$	$\sigma \notin \bar{\sigma}$	(Mod)
$\vdash \text{At}(\bar{F})(\heartsuit) \equiv \heartsuit$	$\heartsuit \in \{\wedge, \neg, \equiv_{\sigma}\}$	(Stab $\heartsuit$)
$\vdash \text{At}(\bar{F})(M) \equiv M$	M a combinator	(Stab λ)
If $\vdash P$ then $\vdash \text{At}(\bar{F})(P)$		(At-Nec)
If $\Gamma \vdash \text{At}(\bar{F})(Q)$ then $\Gamma \vdash \text{At}(\bar{F})(\forall x Q)$	x not free in $\Gamma, \bar{F}$	(At-UG)

FIGURE 3. Domain specifiers.

$\vdash \blacksquare P \rightarrow \text{At}(\bar{F})(P)$	(Master)
$\vdash \blacksquare(M \equiv N) \rightarrow (P[M/x] \rightarrow P[N/x])$	(SLL $\blacksquare$)
If $\Gamma \vdash \text{At}(\bar{F})(\text{At}(\bar{X})(P))$ for all $\bar{\sigma}, \bar{X} : \bar{\sigma}$ with $\bar{X} \notin \text{FV}(\Gamma, P, \bar{F})$, then $\Gamma \vdash \text{At}(\bar{F})(\blacksquare P)$	(Surj)

FIGURE 4. Domain invariability.

A.3. Quantificational Substitution Structures.

Definition A.1. An *applicative structure* is a pair (A, App) , where A is a typed family of sets A^{σ} and App a typed family of functions $\text{App}^{\sigma\tau} : A^{\sigma \rightarrow \tau} \times A^{\sigma} \rightarrow A^{\tau}$. A *substitution structure* is a triple $((I, \circ), (A, \text{App}), \text{Sub})$ where $(I, \circ)$ is a monoid with identity 1, (A, App) is an applicative structure, and Sub is a typed family of maps $I \times A^{\sigma} \rightarrow A^{\sigma}$, written $(i, \mathbf{a}) \mapsto i\mathbf{a}$, such that:

- (21) $1\mathbf{a} = \mathbf{a}$;
- (22) $i(j\mathbf{a}) = (i \circ j)\mathbf{a}$;
- (23) $i \text{App}(\mathbf{f}, \mathbf{a}) = \text{App}(i\mathbf{f}, i\mathbf{a})$.

An element $\mathbf{a}$ is *stable* when $i\mathbf{a} = \mathbf{a}$ for all $i \in I$.

Definition A.2. Let (A, App) be an applicative structure. A *domain* is a partial map $\bar{\mathbf{f}}$ defined on finitely many types, with $\bar{\mathbf{f}}(\sigma) \in A^{\sigma \rightarrow t}$ whenever defined; Dom is the set of domains. A *quantificational substitution structure* (QSS) is a substitution structure equipped with a surjection $\bar{\mathbf{f}} \mapsto i_{\bar{\mathbf{f}}}$ from Dom onto I such that:

- (24) for each type σ there is some $\mathbf{e}_{\sigma} \in A^{\sigma \rightarrow t}$ whose unit domain determines the identity: $i_{\mathbf{e}_{\sigma}} = 1$;
- (25) (Composition) $i_{\bar{\mathbf{f}}} \circ i_{\bar{\mathbf{g}}} = i_{\bar{\mathbf{f}} \bullet \bar{\mathbf{g}}}$ for all domains $\bar{\mathbf{f}}, \bar{\mathbf{g}}$;
- (26) (Identity) $\bar{\mathbf{e}} \bullet \bar{\mathbf{f}} = \bar{\mathbf{f}} \bullet \bar{\mathbf{e}}$ for every domain $\bar{\mathbf{f}}$ and every domain $\bar{\mathbf{e}}$ with $i_{\bar{\mathbf{e}}} = 1$.

Here $(\bar{\mathbf{f}} \bullet \bar{\mathbf{g}})(\rho) = i_{\bar{\mathbf{f}}} \bar{\mathbf{g}}(\rho)$ if $\bar{\mathbf{g}}(\rho)$ is defined, $\bar{\mathbf{f}}(\rho)$ if only $\bar{\mathbf{f}}(\rho)$ is defined, and undefined otherwise.

Given the other conditions, (Identity) is equivalent to requiring that $i_{\bar{\mathbf{f}}}\mathbf{e}_\sigma = \bar{\mathbf{f}}(\sigma)$ whenever $\bar{\mathbf{f}}(\sigma)$ is defined, and $i_{\bar{\mathbf{f}}}\mathbf{e}_\sigma = \mathbf{e}_\sigma$ otherwise. These are the semantic counterparts of (Gen) and (Mod).

A.4. Models. A variable assignment g maps each $x \in \text{Var}^\sigma$ to an element of A^σ . For a nice term sequence $\bar{F}$, write $\llbracket \bar{F} \rrbracket^g$ for the domain mapping σ to $\llbracket F_\sigma \rrbracket^g$ for each $\sigma \in \bar{\sigma}$.

Definition A.3. A $\mathcal{L}_\dagger$ -model consists of a QSS $\mathfrak{A}$, a map $\llbracket \cdot \rrbracket^{(\cdot)}$ taking each $\mathcal{L}_\dagger$ -term $M : \sigma$ and assignment g to an element of A^σ , and a set $\mathcal{T} \subseteq A^t$ of truths, where we write $\mathfrak{A} \models \mathbf{a}$ for $\mathbf{a} \in \mathcal{T}$, such that the following hold for every assignment g and every domain $\bar{\mathbf{f}}$:

- (27) $\llbracket x \rrbracket^g = g(x)$; $\llbracket MN \rrbracket^g = \text{App}(\llbracket M \rrbracket^g, \llbracket N \rrbracket^g)$; $\llbracket M \rrbracket^g = \llbracket N \rrbracket^h$ whenever M, N are $\beta\eta$ -equivalent and g, h agree on their common free variables;
- (28) $i_{\llbracket \mathbf{E}!_\sigma \rrbracket^g} = 1$, and $\llbracket \text{At}(\bar{F})(M) \rrbracket^g = i_{\llbracket \bar{F} \rrbracket^g} \llbracket M \rrbracket^g$;
- (29) $\llbracket \wedge \rrbracket^g, \llbracket \neg \rrbracket^g, \llbracket \equiv \rrbracket^g$ and the denotations of all combinators are stable;
- (30) $\mathfrak{A} \models \llbracket P \wedge Q \rrbracket^g$ iff $\mathfrak{A} \models \llbracket P \rrbracket^g$ and $\mathfrak{A} \models \llbracket Q \rrbracket^g$; $\mathfrak{A} \models \llbracket \neg P \rrbracket^g$ iff $\mathfrak{A} \not\models \llbracket P \rrbracket^g$;
- (31) $\mathfrak{A} \models \llbracket M \equiv N \rrbracket^g$ iff for every $\mathbf{h} \in A^{\sigma \rightarrow t}$, $\mathfrak{A} \models \text{App}(\mathbf{h}, \llbracket M \rrbracket^g)$ iff $\mathfrak{A} \models \text{App}(\mathbf{h}, \llbracket N \rrbracket^g)$;
- (32) $\mathfrak{A} \models i_{\bar{\mathbf{f}}} \llbracket \forall x P \rrbracket^g$ iff $\mathfrak{A} \models i_{\bar{\mathbf{f}}} \llbracket P \rrbracket^{g[x \mapsto \mathbf{a}]}$ for every $\mathbf{a} \in A^\sigma$ with $\mathfrak{A} \models \text{App}(i_{\bar{\mathbf{f}}} \llbracket \mathbf{E}!_\sigma \rrbracket^g, i_{\bar{\mathbf{f}}} \mathbf{a})$, and dually for $\exists$;
- (33) $\mathfrak{A} \models i_{\bar{\mathbf{f}}} \llbracket \blacksquare P \rrbracket^g$ iff $\mathfrak{A} \models (i_{\bar{\mathbf{f}}} \circ i_{\bar{\mathbf{h}}}) \llbracket P \rrbracket^g$ for every $\bar{\mathbf{h}} \in \text{Dom}$.

A formula P is *valid* in a model if $\mathfrak{A} \models \llbracket P \rrbracket^g$ for every assignment g , and *valid simpliciter* if valid in every $\mathcal{L}_\dagger$ -model.

Clauses 6 and 7 build substitutions into the satisfaction conditions because quantifiers and $\blacksquare$, being syncategorematic, receive no independent denotations. The quantificational clause says that the standard free-logical satisfaction conditions hold not only at the current domain but after every quantificational substitution; the clause for $\blacksquare$ treats it as a generalization over all such substitutions.

A.5. Completeness.

Theorem A.4 (Soundness and completeness). *For every $\mathcal{L}_\dagger$ -formula P : $P \in \mathbf{Q}_\dagger$ if and only if P is valid in every $\mathcal{L}_\dagger$ -model.*

Corollary A.5. *$\mathbf{Q}_\dagger$ is consistent, and so is $\mathbf{Q}_\dagger$ extended with quantificationalism.*

The first part of the corollary follows from the existence of $\mathcal{L}_\dagger$ -models, which can be obtained, for example, by taking suitable direct powers of models of classical higher-order logic. For the second, $\mathbf{Q}_\dagger$ proves $\text{At}(\perp_{t \rightarrow t})(\neg \exists p p)$ using (Gen) and the stability axioms: at the empty domain of propositions, $\top$ does not exist. So $\mathbf{E}! \top \wedge \mathbf{E}!(\mathbf{E}! \top)$ implies $\mathbf{E}! \top \wedge \mathbf{E}!(\mathbf{E}! \top) \wedge \blacklozenge \neg \mathbf{E}! \top$, and hence $\exists p(p \wedge \blacklozenge \neg p)$ by (FrUI). Since $\blacksquare$ satisfies T, this yields quantificationalism. As $\mathbf{E}! \top \wedge \mathbf{E}!(\mathbf{E}! \top)$ is satisfiable in some $\mathcal{L}_\dagger$ -model, quantificationalism is consistent with $\mathbf{Q}_\dagger$.